

IEOR 242 - ICO success prediction project

Annotated code

December 22, 2019

0.1 Annotated R Code

0.1.1 Dead coin prediction

```
In [2]: library(caret)
library(MASS)
library(caTools)
library(randomForest)
library(ggplot2)
library(GGally)
library(car)
library(rpart)
library(rattle)
library(boot)
library(dplyr)
library(ROCR)

mean_squared_error <- function(responses, predictions) {
  MSE <- mean(((responses - predictions))^2)
  return(MSE)
}

mean_absolute_error <- function(responses, predictions) {
  MAE <- mean(abs(responses - predictions))
  return(MAE)
}

OS_R_squared <- function(responses, predictions, train_responses) {
  baseline <- mean(train_responses)
  SSE <- sum((responses - predictions)^2)
  SST <- sum((responses - baseline)^2)
  r2 <- 1 - SSE/SST
  return(r2)
}

all_metrics <- function(responses, predictions, train_responses) {
  filter_vec = !is.na(responses) & !is.na(predictions)
```

```

    responses <- responses[filter_vec]
    predictions <- predictions[filter_vec]
    train_responses <- train_responses[filter_vec]
    mse <- mean_squared_error(responses, predictions)
    mae <- mean_absolute_error(responses, predictions)
    OSR2 <- OS_R_squared(responses, predictions, train_responses)
    return(c(mse, mae, OSR2))
}

tableAccuracy <- function(label, pred) {
  t = table(label, pred)
  a = sum(diag(t))/length(label)
  return(a)
}

tableTPR <- function(label, pred) {
  t = table(label, pred)
  return(t[2,2]/(t[2,1] + t[2,2]))
}

tableFPR <- function(label, pred) {
  t = table(label, pred)
  return(t[1,2]/(t[1,1] + t[1,2]))
}

```

We load the preprocessed data:

```

In [3]: df <- read.csv("../python/coincheckup/dataset.csv", stringsAsFactors=FALSE, na="")
df$label_disappeared = as.factor(df$label_disappeared)
# name and Symbol are the only strings

```

a) Logistic regression

We begin with a logistic regression on the features which are numeric and have a very low rate of NAs:

```

In [4]: set.seed(42)
split = sample.split(df$label_disappeared, SplitRatio = 0.7)
train_data <- filter(df, split== TRUE)
test_data <- filter(df, split== FALSE)

table(train_data$label_disappeared)
table(test_data$label_disappeared)

#LOGISTIC
logistic <- glm(label_disappeared ~ Price + X1h + X24h + X7d + X14d + X30d + X45d +
  X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
  Team + Product + Coin + Social + Communication + Business +
  Avg..volume + Age..mo., data = train_data, family="binomial")

```

```
summary(logistic)

pred = predict(logistic, newdata = test_data, type="response")

table(test_data$label_disappeared, pred>0.2)

tableAccuracy(test_data$label_disappeared, pred>0.2)
tableTPR(test_data$label_disappeared, pred>0.2)
tableFPR(test_data$label_disappeared, pred>0.2)

vif(logistic)
```

```
FALSE TRUE
1025   208
```

```
FALSE TRUE
439    89
```

Warning message:

“glm.fit: fitted probabilities numerically 0 or 1 occurred”

Call:

```
glm(formula = label_disappeared ~ Price + X1h + X24h + X7d +
      X14d + X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol +
      Circ..Supply + Total.Supply + Team + Product + Coin + Social +
      Communication + Business + Avg..volume + Age..mo., family = "binomial",
      data = train_data)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.4129	-0.6457	-0.4438	-0.2806	2.6774

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.904e-02	4.592e-01	0.107	0.91494
Price	-5.873e-03	1.089e-02	-0.539	0.58957
X1h	7.936e-05	5.658e-03	0.014	0.98881
X24h	4.684e-04	2.211e-03	0.212	0.83221
X7d	2.764e-03	1.811e-03	1.526	0.12690
X14d	-5.093e-04	1.361e-03	-0.374	0.70821
X30d	5.891e-04	2.114e-03	0.279	0.78049
X45d	-6.051e-03	2.321e-03	-2.608	0.00912 **
X90d	2.879e-03	1.871e-03	1.539	0.12384
X200d	-2.364e-03	2.256e-03	-1.048	0.29464

Mkt..Cap	-5.270e-05	3.726e-04	-0.141	0.88753
X24h.Vol	-4.836e-04	4.349e-04	-1.112	0.26623
Circ..Supply	5.619e-05	4.204e-04	0.134	0.89366
Total.Supply	2.683e-14	4.072e-14	0.659	0.51002
Team	-2.397e-02	1.576e-02	-1.520	0.12847
Product	-4.131e-03	4.869e-03	-0.848	0.39625
Coin	7.949e-03	5.937e-03	1.339	0.18066
Social	-1.858e-02	8.662e-03	-2.145	0.03199 *
Communication	-3.273e-03	3.600e-03	-0.909	0.36335
Business	-3.909e-04	3.243e-03	-0.121	0.90407
Avg..volume	-8.158e-11	8.827e-10	-0.092	0.92636
Age..mo.	7.262e-03	5.277e-03	1.376	0.16878

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1087.97 on 1191 degrees of freedom
 Residual deviance: 954.15 on 1170 degrees of freedom
 (41 observations deleted due to missingness)
 AIC: 998.15

Number of Fisher Scoring iterations: 11

	FALSE	TRUE
FALSE	307	117
TRUE	31	55

0.685606060606061
 0.63953488372093
 0.275943396226415

Price 1.02584781560209 **X1h** 1.15769836383294 **X24h** 2.96033449373743 **X7d** 2.95185035557887
X14d 3.24090275717362 **X30d** 4.91817668782566 **X45d** 4.59783289757329 **X90d** 2.61198284261179
X200d 1.91331591388318 **Mkt..Cap** 1.03647241353168 **X24h.Vol** 1.03491142934378 **Circ..Supply**
 1.05018223690908 **Total.Supply** 1.03860840602233 **Team** 7.10648538827917 **Product**
 1.81658646390026 **Coin** 2.19831931567567 **Social** 3.19603310898613 **Communication**
 1.46875216267816 **Business** 1.5922862721872 **Avg..volume** 1.0164595410715 **Age..mo.**
 1.3693268345375

Then we reduce the range of features to the relevant ones and make a final logistic model, that we compare to the baseline:

```
In [5]: logistic <- glm(label_disappeared ~ X45d + Social + 0, data = train_data,
family="binomial")
summary(logistic)
```

```

pred = predict(logistic, newdata = test_data, type="response")

table(test_data$label_disappeared, pred>0.2)

tableAccuracy(test_data$label_disappeared, pred>0.2)
tableTPR(test_data$label_disappeared, pred>0.2)
tableFPR(test_data$label_disappeared, pred>0.2)

# Baseline accuracy:
t_ <- table(test_data$label_disappeared)
t_[1]/sum(t_)

rocr.pred <- prediction(pred, test_data$label_disappeared)
perf <- performance(rocr.pred, "tpr", "fpr")
plot(perf, colorize = TRUE)
abline(0, 1)
as.numeric(performance(rocr.pred, "auc")@y.values)

```

Warning message:

“glm.fit: fitted probabilities numerically 0 or 1 occurred”

Call:

```

glm(formula = label_disappeared ~ X45d + Social + 0, family = "binomial",
    data = train_data)

```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.1773	-0.6347	-0.4854	-0.3496	2.7011

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
X45d	-0.004062	0.001108	-3.667	0.000245 ***
Social	-0.032039	0.001607	-19.935	< 2e-16 ***

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1


(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1696.8 on 1224 degrees of freedom
 Residual deviance: 1011.9 on 1222 degrees of freedom
 (9 observations deleted due to missingness)
 AIC: 1015.9

Number of Fisher Scoring iterations: 8

	FALSE	TRUE
FALSE	320	117
TRUE	37	52

0.704545454545455
 0.584269662921348
 0.267734553775744
FALSE: 0.831439393939394
 0.708919342812334



b) Random Forest

Let's begin checking the number of NAs in each column to select the features.

```
In [10]: sapply(df, function(x){ sum(is.na(x)) })
```

```
Name 0 MC.. 0 Symbol 0 Price 0 BTC 0 X1h 0 X24h 0 X7d 0 X14d 0 X30d 4 X45d 11 X90d 22
X200d 48 Mkt..Cap 0 MCAP.BTC 0 X24h.Vol 0 X24h.Vol.BTC 0 Circ..Supply 0 Total.Supply 9
Max..Supply 1394 Team 0 Advisors 566 Brand.Buzz 566 Product 0 Coin 0 Social 0
Communication 0 Business 0 GitHub 697 GitHub.1 697 Avg..volume 0 Age..mo. 1
Winning.months 1 label_Price 297 label_Mkt..Cap 297 label_growth_rate_Price 297
label_growth_rate_Mkt..Cap 297 label_disappeared 0
```

Now let's try a random forest with cross-validation. The dataset is not balanced so we try a cross-validation measured with Accuracy and then F1-score as Accuracy is less relevant in this case.

```
In [11]: set.seed(42)
test_data_filled_with_0 <- test_data
test_data_filled_with_0[is.na(test_data_filled_with_0)] <- 0
train_data.mm = as.data.frame(model.matrix(label_disappeared ~ X24h + X14d + X45d +
      Social + Mkt..Cap + Age..mo. + Business, data = train_data))
test_data.mm = as.data.frame(model.matrix(label_disappeared ~ X24h + X14d + X45d +
      Social + Mkt..Cap + Age..mo. + Business, data = test_data_filled_with_0))

train.rf <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
      Age..mo. + Business,
      data = train_data,
      method = "rf",
      na.action = na.omit,
      tuneGrid = data.frame(mtry=1:7),
      trControl = trainControl(method="cv", number=5),
      metric = "Accuracy")

train.rf$results
train.rf
best.rf <- train.rf$finalModel
pred.rf <- predict(best.rf, newdata = test_data.mm) # can use same model matrix

ggplot(train.rf$results, aes(x = mtry, y = Accuracy)) + geom_point(size = 3) +
  ylab("CV Accuracy") + theme_bw() + theme(axis.title=element_text(size=18),
      axis.text=element_text(size=18))
```

mtry <int>	Accuracy <dbl>	Kappa <dbl>	AccuracySD <dbl>	KappaSD <dbl>
1	0.8480651	0.1820981	0.01230487	0.07219353
2	0.8497045	0.2675792	0.01730529	0.08837830
3	0.8545991	0.3131013	0.01457905	0.07484737
4	0.8505074	0.2927120	0.01216925	0.05062367
5	0.8505175	0.3181606	0.01943042	0.07520890
6	0.8472421	0.2995333	0.01863631	0.06998712
7	0.8407181	0.2733978	0.02110235	0.08098173

Random Forest

1233 samples

7 predictor

2 classes: 'FALSE', 'TRUE'

No pre-processing

Resampling: Cross-Validated (5 fold)


Summary of sample sizes: 980, 978, 980, 979, 979

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
1	0.8480651	0.1820981
2	0.8497045	0.2675792
3	0.8545991	0.3131013
4	0.8505074	0.2927120
5	0.8505175	0.3181606
6	0.8472421	0.2995333
7	0.8407181	0.2733978

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was mtry = 3.



```
In [14]: f1 <- function (data, lev = NULL, model = NULL) {
  pred <- data$pred[!is.na(data$pred)&!is.na(data$obs)]
  obs <- data$obs[!is.na(data$pred)&!is.na(data$obs)]
  precision <- posPredValue(pred, obs, positive = lev[2])
  recall <- sensitivity(pred, obs, postive = lev[1])
  f1_val <- (2 * precision * recall) / (precision + recall)
  names(f1_val) <- c("F1")
  #print(precision)
  #print(recall)
  #print(f1_val)
  f1_val
}
```

```

train.rf <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business,
  data = train_data,
  method = "rf",
  na.action = na.omit,
  tuneGrid = data.frame(mtry=1:7),
  trControl = trainControl(method="cv", number=5, summaryFunction = f1),
  metric = PY1+s"F1")

train.rf$results
train.rf
best.rf <- train.rf$finalModel
pred.rf <- predict(best.rf, newdata = test_data.mm) # can use same model matrix

ggplot(train.rf$results, aes(x = mtry, y = F1)) + geom_point(size = 3) +
  ylab("CV F1") + theme_bw() + theme(axis.title=element_text(size=18),
  axis.text=element_text(size=18))

```

mtry <int>	F1 <dbl>	F1SD <dbl>
1	0.9028305	0.08709772
2	0.7850726	0.07805387
3	0.7611599	0.11121317
4	0.7696718	0.06996851
5	0.7266709	0.10382716
6	0.7246719	0.11886381
7	0.7067876	0.12034424


Random Forest

1233 samples
 7 predictor
 2 classes: 'FALSE', 'TRUE'

No pre-processing
 Resampling: Cross-Validated (5 fold)
 Summary of sample sizes: 979, 980, 979, 978, 980
 Resampling results across tuning parameters:

mtry	F1
1	0.9028305
2	0.7850726
3	0.7611599
4	0.7696718
5	0.7266709
6	0.7246719
7	0.7067876

F1 was used to select the optimal model using the largest value.
 The final value used for the model was mtry = 1.



An mtry of 3 looks good. We can try sampling stratification to handle the imbalance of the data.

```
In [15]: set.seed(42)
rare.class.prevalence = 0.2
nRareSamples = 1000 * rare.class.prevalence
mod.rf <- randomForest(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business, data = train_data, mtry = 3, nodesize = 5, ntree = 1000,
  strata=train_data$label_disappeared,
  sampsize=c(nRareSamples,nRareSamples), na.action = na.omit)
# print(mod.rf)
```

```

pred.rf <- predict(mod.rf, newdata = test_data)

table(test_data$label_disappeared, pred.rf)

tableAccuracy(test_data$label_disappeared, pred.rf)
tableTPR(test_data$label_disappeared, pred.rf)
tableFPR(test_data$label_disappeared, pred.rf)

pred.rf
  FALSE TRUE
FALSE  427  10
TRUE   70  19

0.84469696969697
0.213483146067416
0.022883295194508

```

c) Boosting

Finally we can try a boosting method with cross-validation.

```

In [17]: tGrid = expand.grid(n.trees = (1:10)*1000, interaction.depth = c(1,2,4,6,8,10),
                             shrinkage = 0.001, n.minobsinnode = 10)

set.seed(42)
train.boost <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
                      Age..mo. + Business,
                      data = train_data,
                      method = "gbm",    ## gradient boosting machine
                      tuneGrid = tGrid,
                      trControl = trainControl(method="cv", number=5),
                      metric = "Accuracy",
                      na.action = na.omit)

train.boost
best.boost <- train.boost$finalModel

ggplot(train.boost$results, aes(x = n.trees, y = Accuracy, colour =
  as.factor(interaction.depth))) + geom_line() +
  ylab("CV Accuracy") + theme_bw() + theme(axis.title=element_text(size=18),
  axis.text=element_text(size=18)) +
  scale_color_discrete(name = "interaction.depth")

```

Stochastic Gradient Boosting

```

1233 samples
  7 predictor
  2 classes: 'FALSE', 'TRUE'

```

No pre-processing

Resampling: Cross-Validated (5 fold)


Summary of sample sizes: 980, 978, 980, 979, 979

Tuning parameter 'shrinkage' was held constant at a value of 0.001

Tuning parameter 'n.minobsinnode' was held constant at a value of 10

Accuracy was used to select the optimal model using the largest value.

The final values used for the model were n.trees = 2000, interaction.depth = 8, shrinkage = 0.001 and n.minobsinnode = 10.



We choose n.trees = 2000, interaction.depth = 8.

```
In [18]: pred.best.boost <- predict(best.boost, newdata = test_data.mm, n.trees = 2000,
```

```
interaction.depth = 8) # from CV plot

table(test_data$label_disappeared, pred.best.boost>0.7)

tableAccuracy(test_data$label_disappeared, pred.best.boost>0.7)
tableTPR(test_data$label_disappeared, pred.best.boost>0.7)
tableFPR(test_data$label_disappeared, pred.best.boost>0.7)
```

	FALSE	TRUE
FALSE	221	218
TRUE	70	19

```
0.454545454545455
0.213483146067416
0.496583143507973
```

0.1.2 Portfolio determination based on naive growth rates

a) Regression

We try a linear model on a lot of features.

```
In [22]: set.seed(42)
split = sample.split(df$label_growth_rate_Price, SplitRatio = 0.7)
price_train_data <- filter(df, split== TRUE)
price_test_data <- filter(df, split== FALSE)

price_train_data = price_train_data %>% filter(as.integer(label_disappeared)==1)
# we keep label_disappeared = FALSE
price_test_data = price_test_data %>% filter(as.integer(label_disappeared)==1)

mod <- lm(label_growth_rate_Price ~ Price + X1h + X24h + X7d + X14d + X30d + X45d +
          X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
          Team + Product + Coin + Social + Communication + Business +
          Avg..volume + Age..mo., data = price_train_data)

summary(mod)

pred_lm = predict(mod, newdata = price_test_data)

all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
            price_train_data$label_growth_rate_Price)
```

Call:

```
lm(formula = label_growth_rate_Price ~ Price + X1h + X24h + X7d +
    X14d + X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol +
```

```
Circ..Supply + Total.Supply + Team + Product + Coin + Social +
Communication + Business + Avg..volume + Age..mo., data = price_train_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-95.323	-0.448	0.006	0.504	157.521

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.380e+00	1.463e+00	-0.943	0.346
Price	-1.363e-05	9.791e-04	-0.014	0.989
X1h	2.563e-02	2.873e-02	0.892	0.373
X24h	-1.096e-02	9.764e-03	-1.122	0.262
X7d	-3.442e-03	7.134e-03	-0.482	0.630
X14d	9.311e-04	3.500e-03	0.266	0.790
X30d	5.279e-04	4.867e-03	0.108	0.914
X45d	-2.804e-03	2.983e-03	-0.940	0.347
X90d	5.591e-03	5.494e-03	1.018	0.309
X200d	-2.616e-03	5.139e-03	-0.509	0.611
Mkt..Cap	4.964e-05	1.062e-03	0.047	0.963
X24h.Vol	9.059e-06	1.117e-03	0.008	0.994
Circ..Supply	-1.920e-03	1.103e-03	-1.741	0.082 .
Total.Supply	9.642e-12	3.189e-13	30.230	<2e-16 ***
Team	1.977e-02	4.204e-02	0.470	0.638
Product	-6.773e-03	1.328e-02	-0.510	0.610
Coin	-8.403e-03	1.685e-02	-0.499	0.618
Social	3.590e-03	2.370e-02	0.151	0.880
Communication	-5.365e-03	8.557e-03	-0.627	0.531
Business	6.493e-03	1.123e-02	0.578	0.563
Avg..volume	2.020e-11	1.844e-09	0.011	0.991
Age..mo.	1.896e-02	1.633e-02	1.161	0.246

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.24 on 962 degrees of freedom

(40 observations deleted due to missingness)

Multiple R-squared: 0.4896, Adjusted R-squared: 0.4785

F-statistic: 43.95 on 21 and 962 DF, p-value: < 2.2e-16

1. 582.295129943916 2. 2.15332619393653 3. -173.530306079791

In [23]: vif(mod)

Price 1.02423943344359 X1h 1.27695002473311 X24h 2.20857264825821 X7d 2.20262213596122
X14d 51.0399053125082 X30d 57.8341308908584 X45d 16.3502969281802 X90d 6.77458705855804
X200d 3.77206853766809 Mkt..Cap 1.02694209743189 X24h.Vol 1.02876014673704 Circ..Supply

```

1.07660848685341 Total.Supply      1.01072167422902 Team      6.49637787342385 Product
1.9710382752697 Coin      1.90590139973695 Social      3.14333047583439 Communication
1.53408207175482 Business      1.36185221776999 Avg..volume      1.02165670821974 Age..mo.
1.3028447664378

```

This is not good, there is some multicollinearity. Let's remove the guilty features.

```

In [26]: mod <- lm(label_growth_rate_Price ~ Price + X1h + X24h + X7d + X30d + X90d + X200d +
      Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team + Product +
      Coin + Social + Communication + Business + Avg..volume + Age..mo.,
      data = price_train_data)

```

```
summary(mod)
```

```
pred_lm = predict(mod, newdata = price_test_data)
```

```
all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
            price_train_data$label_growth_rate_Price)
```

Call:

```

lm(formula = label_growth_rate_Price ~ Price + X1h + X24h + X7d +
    X30d + X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply +
    Total.Supply + Team + Product + Coin + Social + Communication +
    Business + Avg..volume + Age..mo., data = price_train_data)

```

Residuals:

Min	1Q	Median	3Q	Max
-95.222	-0.426	-0.009	0.496	157.685

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.506e+00	1.452e+00	-1.037	0.2999
Price	-1.311e-05	9.783e-04	-0.013	0.9893
X1h	2.299e-02	2.857e-02	0.805	0.4212
X24h	-9.303e-03	9.536e-03	-0.976	0.3295
X7d	-3.687e-03	7.124e-03	-0.517	0.6049
X30d	-8.050e-05	9.549e-04	-0.084	0.9328
X90d	1.554e-03	3.466e-03	0.448	0.6541
X200d	-4.926e-04	4.608e-03	-0.107	0.9149
Mkt..Cap	5.801e-05	1.061e-03	0.055	0.9564
X24h.Vol	7.150e-05	1.114e-03	0.064	0.9489
Circ..Supply	-1.902e-03	1.102e-03	-1.726	0.0846 .
Total.Supply	9.631e-12	3.186e-13	30.233	<2e-16 ***
Team	1.702e-02	4.190e-02	0.406	0.6847
Product	-6.461e-03	1.327e-02	-0.487	0.6263
Coin	-7.157e-03	1.676e-02	-0.427	0.6695
Social	5.232e-03	2.359e-02	0.222	0.8245
Communication	-5.266e-03	8.549e-03	-0.616	0.5381


```

Business      6.734e-03  1.122e-02  0.600  0.5484
Avg..volume   1.029e-11  1.843e-09  0.006  0.9955
Age..mo.      2.052e-02  1.613e-02  1.272  0.2037
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

Residual standard error: 7.236 on 964 degrees of freedom
(40 observations deleted due to missingness)
Multiple R-squared:  0.4891, Adjusted R-squared:  0.4791
F-statistic: 48.58 on 19 and 964 DF,  p-value: < 2.2e-16

```

```
1. 564.062942832866 2. 1.90025594030086 3. -168.065604361777
```

```
In [27]: vif(mod)
```

```

Price 1.02383458947674 X1h 1.26378621144891 X24h 2.10915904620629 X7d 2.1991343191775
X30d 2.22872902583279 X90d 2.69968306031107 X200d 3.03647959935625 Mkt..Cap
1.02665757634139 X24h.Vol 1.02508903973853 Circ..Supply 1.07610953486792 Total.Supply
1.00943241102529 Team 6.46108652227192 Product 1.96792266463445 Coin 1.88882024934975
Social 3.11853701685625 Communication 1.53321255914753 Business 1.36088711654353
Avg..volume 1.02158918501001 Age..mo. 1.27210148432654

```

This still does not look good. Let's narrow down the features.

```
In [20]: mod <- lm(label_growth_rate_Price ~ Circ..Supply + Total.Supply + 0,
                  data = price_train_data)
```

```
summary(mod)
```

```
pred_lm = predict(mod, newdata = price_test_data)
```

```
all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
            price_train_data$label_growth_rate_Price)
```

Call:

```
lm(formula = label_growth_rate_Price ~ Circ..Supply + Total.Supply +
    0, data = price_train_data)
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-96.415  -0.701  -0.407   0.164 158.564

```

Coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
Circ..Supply -2.534e-03  8.533e-04  -2.97  0.00305 **
Total.Supply  9.579e-12  3.097e-13  30.93 < 2e-16 ***

```

```
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.074 on 1017 degrees of freedom

(5 observations deleted due to missingness)

Multiple R-squared: 0.485, Adjusted R-squared: 0.484

F-statistic: 478.9 on 2 and 1017 DF, p-value: < 2.2e-16

1. 547.59691212069 2. 1.94101815834009 3. -166.769645220578

We can compare it to the baseline (predicting negative all the time). We assume that we invest the same amount of money in each coin, ignoring fees, and compute the result growth rate of our portfolio following our regression strategy and the baseline.

```
In [29]: filter_vec = !is.na(pred_lm)
          sum((pred_lm*price_test_data$label_growth_rate_Price)[filter_vec])
          /sum(pred_lm[filter_vec])

          # Baseline
          sum(price_test_data$label_growth_rate_Price[filter_vec])
          /length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.142130885876864

-0.342125271527515

We beat it!

0.1.3 Portfolio determination in a log normal stock model

a) Regression

With our new log-normal assumption we can try to predict the difference of the logarithms of the price between March and December.

```
In [30]: df$price_log_diff = log(df$label_Price) - log(df$Price) # follow  $\mu T + \sigma B_T$ 

          split = sample.split(df$price_log_diff, SplitRatio = 0.7)
          price_train_data <- filter(df, split== TRUE)
          price_test_data <- filter(df, split== FALSE)

          price_train_data = price_train_data %>% filter(as.integer(label_disappeared)==1)
          # we keep label_disappeared = FALSE
          price_test_data = price_test_data %>% filter(as.integer(label_disappeared)==1)

          mod_log <- lm(price_log_diff ~ Price + X1h + X24h + X7d + X14d + X30d + X45d + X90d +
                        X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team +
                        Product + Coin + Social + Communication + Business + Avg..volume + Age..mo.,
                        data = price_train_data)

          summary(mod_log)
```

```

pred_lm = predict(mod_log, newdata = price_test_data)

all_metrics(price_test_data$price_log_diff, pred_lm, price_train_data$price_log_diff)

Call:
lm(formula = price_log_diff ~ Price + X1h + X24h + X7d + X14d +
    X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply +
    Total.Supply + Team + Product + Coin + Social + Communication +
    Business + Avg..volume + Age..mo., data = price_train_data)

Residuals:
    Min       1Q   Median       3Q      Max
-6.2257 -0.4446  0.0628  0.5709  3.8110

Coefficients:
                Estimate Std. Error t value Pr(>|t|)
(Intercept)   -9.089e-01  1.915e-01  -4.747 2.37e-06 ***
Price          2.427e-05  3.484e-05   0.697  0.4861
X1h           -3.607e-03  5.843e-03  -0.617  0.5372
X24h           -3.514e-03  1.624e-03  -2.164  0.0307 *
X7d            -1.298e-03  1.028e-03  -1.262  0.2074
X14d           -6.971e-04  3.174e-04  -2.196  0.0283 *
X30d            5.989e-04  4.839e-04   1.238  0.2161
X45d            6.502e-05  2.585e-04   0.252  0.8014
X90d            5.886e-04  5.132e-04   1.147  0.2517
X200d          -9.859e-04  5.229e-04  -1.885  0.0597 .
Mkt..Cap       -2.969e-04  1.388e-04  -2.139  0.0327 *
X24h.Vol       -7.968e-07  1.410e-04  -0.006  0.9955
Circ..Supply    2.499e-04  1.455e-04   1.718  0.0862 .
Total.Supply    1.668e-14  1.874e-14   0.890  0.3738
Team            7.936e-03  5.567e-03   1.426  0.1543
Product        -1.076e-03  1.750e-03  -0.615  0.5385
Coin           -2.972e-03  2.195e-03  -1.354  0.1761
Social          -7.665e-04  3.138e-03  -0.244  0.8071
Communication    3.429e-04  1.144e-03   0.300  0.7645
Business       -3.736e-03  1.468e-03  -2.544  0.0111 *
Avg..volume     1.576e-10  1.418e-10   1.112  0.2665
Age..mo.        3.241e-03  2.212e-03   1.465  0.1433
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9755 on 978 degrees of freedom
(37 observations deleted due to missingness)
Multiple R-squared:  0.05583, Adjusted R-squared:  0.03555
F-statistic: 2.754 on 21 and 978 DF,  p-value: 3.95e-05

```

1. 1.37419565150063 2. 0.786600323768397 3. -0.0510270204478771

```
In [31]: vif(mod_log)
```

```
Price 1.06407441227154 X1h 1.07115978177007 X24h 1.22544188241185 X7d 1.31330264496053
X14d 753.264205918813 X30d 1730.32601194076 X45d 491.069203199915 X90d 10.6815410522711
X200d 3.59844428634785 Mkt..Cap 1.0247473315568 X24h.Vol 1.02517852100103 Circ..Supply
1.07871229383061 Total.Supply 1.01891666805316 Team 6.3719217947138 Product
1.92537905891746 Coin 1.86682147289798 Social 3.14543604788404 Communication
1.52506691080304 Business 1.3559062673666 Avg..volume 1.09292020780236 Age..mo.
1.31230648270897
```

```
In [34]: mod_log <- lm(price_log_diff ~ Price + X1h + X24h + X7d + X14d + X90d + X200d +
Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team + Product +
Coin + Social + Communication + Business + Avg..volume + Age..mo.,
data = price_train_data)
```

```
summary(mod_log)
```

```
pred_lm = predict(mod_log, newdata = price_test_data)
```

```
all_metrics(price_test_data$price_log_diff, pred_lm, price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ Price + X1h + X24h + X7d + X14d +
X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
Team + Product + Coin + Social + Communication + Business +
Avg..volume + Age..mo., data = price_train_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.1966	-0.4454	0.0553	0.5753	3.7657

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-9.555e-01	1.909e-01	-5.005	6.62e-07	***
Price	2.525e-05	3.491e-05	0.723	0.469718	
X1h	-3.627e-03	5.851e-03	-0.620	0.535459	
X24h	-3.597e-03	1.626e-03	-2.212	0.027225	*
X7d	-1.328e-03	1.029e-03	-1.291	0.196887	
X14d	-4.881e-05	2.479e-05	-1.969	0.049266	*
X90d	8.115e-04	4.038e-04	2.010	0.044752	*
X200d	-1.469e-03	4.261e-04	-3.448	0.000589	***
Mkt..Cap	-2.973e-04	1.391e-04	-2.138	0.032777	*
X24h.Vol	-6.626e-06	1.412e-04	-0.047	0.962579	
Circ..Supply	2.468e-04	1.458e-04	1.693	0.090731	.

```

Total.Supply  1.546e-14  1.874e-14  0.825 0.409431
Team          8.054e-03  5.568e-03  1.446 0.148377
Product       -1.288e-03  1.751e-03  -0.735 0.462238
Coin          -2.875e-03  2.198e-03  -1.308 0.191143
Social        -5.695e-04  3.138e-03  -0.181 0.856039
Communication  4.538e-04  1.145e-03  0.396 0.692061
Business      -3.784e-03  1.471e-03  -2.573 0.010242 *
Avg..volume   1.605e-10  1.420e-10  1.130 0.258656
Age..mo.      3.300e-03  2.204e-03  1.497 0.134725
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

Residual standard error: 0.9776 on 980 degrees of freedom
(37 observations deleted due to missingness)
Multiple R-squared:  0.04984, Adjusted R-squared:  0.03142
F-statistic: 2.706 on 19 and 980 DF,  p-value: 0.0001092

```

```
1. 1.37718128083273 2. 0.789451576546088 3. -0.0533105214163541
```

```
In [35]: vif(mod_log)
```

```

Price 1.06389350752336 X1h 1.06954219074384 X24h 1.22325176284725 X7d 1.30815268834064
X14d  4.57576438585392 X90d  6.58402089382897 X200d  2.37953671288416 Mkt..Cap
1.02437219200182 X24h.Vol  1.02414192975333 Circ..Supply  1.07809508758456 Total.Supply
1.01415980805924 Team  6.34760662656953 Product  1.92036647546332 Coin  1.86268778843273
Social  3.13312064680212 Communication  1.52216803443362 Business  1.3550512673174
Avg..volume  1.09262600492207 Age..mo.  1.29729721274557

```

The VIF is okay now. Let's remove some features.

```
In [38]: mod_log <- lm(price_log_diff ~ X24h + X200d + Mkt..Cap + Business,
                        data = price_train_data)
```

```
summary(mod_log)
```

```
pred_lm_log = predict(mod_log, newdata = price_test_data)
```

```
all_metrics(price_test_data$price_log_diff, pred_lm_log,
            price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ X24h + X200d + Mkt..Cap + Business,
    data = price_train_data)
```

Residuals:

```
Min      1Q  Median      3Q      Max
```

-6.2460 -0.4330 0.0580 0.5923 3.9618

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.7949174	0.1167271	-6.810	1.68e-11	***
X24h	-0.0047998	0.0014715	-3.262	0.00114	**
X200d	-0.0008992	0.0002755	-3.265	0.00113	**
Mkt..Cap	-0.0003394	0.0001378	-2.464	0.01392	*
Business	-0.0027869	0.0012598	-2.212	0.02717	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9794 on 1000 degrees of freedom

(32 observations deleted due to missingness)

Multiple R-squared: 0.03057, Adjusted R-squared: 0.02669

F-statistic: 7.883 on 4 and 1000 DF, p-value: 2.956e-06

1. 1.35643667936919 2. 0.784982884758303 3. -0.0401929076627052

```
In [39]: filter_vec = !is.na(pred_lm_log)
sum((exp(pred_lm_log)*price_test_data$label_growth_rate_Price)[filter_vec])
/sum(exp(pred_lm_log)[filter_vec])

# Baseline
sum(price_test_data$label_growth_rate_Price[filter_vec])
/length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.461436810190696

0.439390893108447

That's fine but not amazing, let's think a bit and try something else.

```
In [40]: mod_log <- lm(price_log_diff ~ X7d + Mkt..Cap + Age..mo., data = price_train_data)

summary(mod_log)

pred_lm_log = predict(mod_log, newdata = price_test_data)

all_metrics(price_test_data$price_log_diff, pred_lm_log,
            price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ X7d + Mkt..Cap + Age..mo., data = price_train_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.3005	-0.4493	0.0632	0.5849	3.9844

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.0571347	0.0505030	-20.932	<2e-16 ***
X7d	-0.0016871	0.0009027	-1.869	0.0619 .
Mkt..Cap	-0.0003295	0.0001397	-2.359	0.0185 *
Age..mo.	0.0041164	0.0019420	2.120	0.0343 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9966 on 1033 degrees of freedom

Multiple R-squared: 0.01299, Adjusted R-squared: 0.01012

F-statistic: 4.531 on 3 and 1033 DF, p-value: 0.003655

1. 1.3321906348193 2. 0.766205419022985 3. 0.0152780706337514

```
In [41]: filter_vec = !is.na(pred_lm_log)
sum((exp(pred_lm_log)*price_test_data$label_growth_rate_Price)[filter_vec])
/sum(exp(pred_lm_log)[filter_vec])

# Baseline
sum(price_test_data$label_growth_rate_Price[filter_vec])
/length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.500360688948723

0.412744498999508

Much better already! We see that filtering out coins where informations are missing, our portfolio behaves already better! That's because trustworthy coins get all their information filled.

b) Random forest

We take the features from the 2 previous models and we try it out on a random forest.

```
In [50]: price_test_data_filled_with_0 <- price_test_data
price_test_data_filled_with_0[is.na(price_test_data_filled_with_0)] <- 0
price_test_data_log.mm = as.data.frame(model.matrix(price_log_diff ~ X24h + X7d +
X200d + Mkt..Cap + Age..mo. + Business, data = price_test_data_filled_with_0))

train.rf <- train(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. +
Business,
data = price_train_data,
method = "rf",
na.action = na.omit,
tuneGrid = data.frame(mtry=1:6),
trControl = trainControl(method="cv", number=5),
distribution="gaussian",
metric = "RMSE")
```

```

train.rf$results
train.rf
best.rf <- train.rf$finalModel
pred.rf <- predict(best.rf, newdata = price_test_data_log.mm)
all_metrics(price_test_data$price_log_diff, pred_lm_log,
             price_train_data$price_log_diff)

```

mtry <int>	RMSE <dbl>	Rsquared <dbl>	MAE <dbl>	RMSESD <dbl>	RsquaredSD <dbl>	MAESD <dbl>
1	0.9868626	0.02396270	0.7021205	0.09754761	0.03461219	0.04148504
2	1.0030523	0.02013034	0.7180743	0.09988699	0.03206543	0.03882143
3	1.0138758	0.01549198	0.7273112	0.10018572	0.02833747	0.03858953
4	1.0172491	0.01483976	0.7298685	0.09965039	0.02697329	0.04055537
5	1.0190250	0.01714539	0.7307562	0.10116483	0.03156220	0.03939837
6	1.0224764	0.01754349	0.7330653	0.10168716	0.03172612	0.04190727

Random Forest

1037 samples
6 predictor

No pre-processing

Resampling: Cross-Validated (5 fold)

Summary of sample sizes: 804, 805, 805, 804, 802

Resampling results across tuning parameters:

mtry	RMSE	Rsquared	MAE
1	0.9868626	0.02396270	0.7021205
2	1.0030523	0.02013034	0.7180743
3	1.0138758	0.01549198	0.7273112
4	1.0172491	0.01483976	0.7298685
5	1.0190250	0.01714539	0.7307562
6	1.0224764	0.01754349	0.7330653

RMSE was used to select the optimal model using the smallest value.


The final value used for the model was mtry = 1.

1. 1.3321906348193 2. 0.766205419022985 3. 0.0152780706337514

```

In [51]: ggplot(train.rf$results, aes(x = mtry, y = Rsquared)) + geom_point(size = 3) +
         ylab("CV Rsquared") + theme_bw() + theme(axis.title=element_text(size=18),
         axis.text=element_text(size=18))

```

The best value is 1 for mtry.

```
In [52]: filter_vec = !is.na(pred.rf)
          sum((exp(pred.rf)*price_test_data$label_growth_rate_Price)[filter_vec])
          /sum(exp(pred.rf)[filter_vec])

          # Baseline
          sum(price_test_data$label_growth_rate_Price[filter_vec])
          /length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.318126401502254

0.410949916909122

Unfortunately this model does not beat the baseline.

c) Boosting

We use the same features.

```
In [54]: tGrid = expand.grid(n.trees = (1:10)*1000, interaction.depth = c(1,2,4,6,8,10),
                             shrinkage = 0.001, n.minobsinnode = 10)

set.seed(42)
train.boost <- train(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. +
                     Business,
                     data = price_train_data,
                     method = "gbm",    ## gradient boosting machine
                     tuneGrid = tGrid,
                     trControl = trainControl(method="cv", number=5),
                     distribution = "gaussian",
                     metric = "RMSE",
                     na.action = na.omit)

train.boost
best.boost <- train.boost$finalModel

ggplot(train.boost$results, aes(x = n.trees, y = Rsquared, colour =
  as.factor(interaction.depth))) + geom_line() +
  ylab("CV Rsquared") + theme_bw() + theme(axis.title=element_text(size=18),
  axis.text=element_text(size=18)) +
  scale_color_discrete(name = "interaction.depth")
```


Stochastic Gradient Boosting

1037 samples
6 predictor

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 803, 804, 805, 804, 804

Tuning parameter 'shrinkage' was held constant at a value of 0.001

Tuning parameter 'n.minobsinnode' was held constant at a value of 10
RMSE was used to select the optimal model using the smallest value.
The final values used for the model were n.trees = 1000, interaction.depth =
6, shrinkage = 0.001 and n.minobsinnode = 10.



The best parameters are the biggest ones.

```
In [56]: pred.best.boost <- predict(best.boost, newdata = price_test_data_log.mm,
  n.trees = 10000, interaction.depth=10) # from CV plot
```

Warning message in predict.gbm(best.boost, newdata = price_test_data_log.mm, n.trees = 10000, :
 “Number of trees not specified or exceeded number fit so far. Using 1000.”

It does not want to do it so we need to do it ourselves.

```
In [63]: library(gbm)

mod.boost <- gbm(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. + Business,
  data = price_train_data,
```

```

distribution = "gaussian",
n.trees = 10000,
shrinkage = 0.001,
interaction.depth = 10)

# NOTE: we need to specify number of trees to get a prediction for boosting
pred.boost <- predict(mod.boost, newdata = price_test_data_log.mm, n.trees=10000)

In [64]: filter_vec = !is.na(pred.boost)
sum((exp(pred.boost)*price_test_data$label_growth_rate_Price)[filter_vec])
      /sum(exp(pred.boost)[filter_vec])

# Baseline
sum(price_test_data$label_growth_rate_Price[filter_vec])
      /length(price_test_data$label_growth_rate_Price[filter_vec])

0.290123603135029
0.410949916909122
This is also very bad compared to the baseline. Too much overfitting!

```