



COLLEGE OF ENGINEERING - Fall 2019

INDUSTRIAL ENGINEERING AND OPERATIONS RESEARCH  
DEPARTMENT - DATA SCIENCE (INDENG242)

COURSE PROJECT:

# Machine Learning For The Prediction Of The Success Of Altcoins

---

Made by:

Johnson Kanjirathinkal, Philippe Ferreira De Sousa,  
Alejandro Sanchez Molina, Arun Putcha



# Contents

<b>1</b>	<b>Motivation</b>	<b>2</b>
<b>2</b>	<b>Data</b>	<b>2</b>
<b>3</b>	<b>Dead coin prediction</b>	<b>4</b>
3.1	Logistic regression . . . . .	4
3.2	Random forest . . . . .	5
3.3	Gradient Boosting . . . . .	6
3.4	Conclusion . . . . .	7
<b>4</b>	<b>Portfolio determination based on naive growth rates</b>	<b>7</b>
<b>5</b>	<b>Portfolio determination in a log normal stock model</b>	<b>8</b>
5.1	Linear regression . . . . .	8
5.2	Random forest . . . . .	8
5.3	Gradient boosting . . . . .	8
5.4	Conclusion . . . . .	9
<b>6</b>	<b>Portfolio determination in real-world</b>	<b>10</b>
<b>7</b>	<b>Appendix</b>	<b>11</b>
7.1	Column Definitions . . . . .	11
7.2	Python Code . . . . .	13
7.3	Annotated R Code . . . . .	21

# 1 Motivation

Altcoins are alternative cryptocurrencies launched after the success of bitcoin. While the success of altcoins primarily depends on a number of factors such as the tech use case, funding, community and competition, general volume movement surges after they get listed on a major exchange.



Figure 1: Value of \$100 of IOTA invested at the time of ICO

With the thousands of Initial Coin Offerings (ICOs), airdrops and obscure offerings that altcoins end up going through, we want to predict the likelihood of a particular set of altcoins gaining in value over a period of time. For instance, Fig. 1 represents the value of 100\$ invested in IOTA at the time of ICO till date. The all time high was equivalent to \$910,170 and \$35,080 at the time of this report. Even if some of the portfolio of altcoins selected lose its value, the large growth rates of a small number would relate to profitability. This calculation has been considered for determination of portfolio value.

Identifying what makes an altcoin successful or unsuccessful has been hypothesized to be dependent on features not limited to the existence of a white paper, team credibility, Github activity, market cap etc. The data collection and modelling has been done to test these hypotheses.

# 2 Data

Historical altcoin data is not easily available publicly as a result of a variety of coins constantly being added and removed. We determined that Coinmarketcap.com and Coincheckup.com provided sufficient resources for data collection pertaining to our problem statement. In order to evaluate the change in altcoin performance, we decided to consider a 10-month gap between our prediction estimation. All required data for this time frame was scrapped from the HTML from a historical snapshot from The Wayback Machine.

The data from Feb 2019 had 1761 observations, while the data from Dec 2019 had 1916 observations. The count of common coins between these data sets was found to be 1460. Fields that were considered to provide no value to the contribution of features was removed. Fields that had a large number of “NAs” were also dropped. As a consequence of altcoins purchases being done in exchange for BTC or ETH, all prices were converted to BTC, tied to the USD value of BTC on

#	Name	MC	Symbol	Price	BTC	1h	24h	7d	14d	30d	45d	90d	200d	Mkt. Cap	MCAP BTC	24h Vol	24h Vol BTC	Circ. Supply	Total Supply
1	Cosmos	16	ATOM	\$ 4.17	0.00057941	1.40%	-1.15%	-5.62%	6.14%	41.73%	11.60%	50.61%	-32.12%	\$ 795.32 MM		\$ 112.64 MM		190.688 MM	237,928,231
2	Qx	48	ZRX	\$ 0.1908	0.00002651	0.33%	-0.51%	-7.61%	18.39%	-18.54%	-37.61%	-19.22%	-39.26%	\$ 115.32 MM		\$ 11.14 MM		604.422 MM	1,000,000,000
3	Monero	14	XMR	\$ 47.14	0.00654831	0.25%	1.79%	-8.35%	-12.27%	-13.79%	-25.79%	-33.53%	-45.02%	\$ 818.55 MM		\$ 110.33 MM		17.365 MM	17,365,457
4	Zilliqa	93	ZIL	\$ 0.0050	0.00000069	0.40%	-2.27%	-3.10%	-28.40%	3.91%	-23.56%	-25.73%	-77.38%	\$ 48.40 MM		\$ 6.74 MM		9.768 MM	13,059,245,347
5	IOTA	25	MIOTA	\$ 0.1622	0.00002253	0.39%	-1.72%	-14.62%	-20.74%	-22.47%	-42.13%	-43.69%	-62.34%	\$ 450.78 MM		\$ 8.63 MM		2.780 MM	2,779,530,283
6	Dash	27	DASH	\$ 43.26	0.00601012	0.85%	-0.18%	-13.34%	-17.63%	-23.47%	-40.60%	-52.55%	-69.87%	\$ 399.01 MM		\$ 204.50 MM		9.223 MM	9,222,896
7	Apollo Currency	260	APL	\$ 0.0007	0.00000010	1.27%	3.89%	3.87%	-11.84%	-13.07%	-52.13%	-36.29%	-44.48%	\$ 10.12 MM		\$ 832.51 K		14.685 Bn	21,165,096,531
8	PIVX	229	PIVX	\$ 0.2209	0.00003069	0.08%	-0.21%	-0.93%	-3.85%	12.09%	-6.55%	-15.89%	-67.23%	\$ 12.54 MM		\$ 152.43 K		56.781 MM	56,781,166
9	Nexo	87	NEXO	\$ 0.0934	0.00001297	0.52%	1.04%	1.12%	-3.95%	5.65%	-9.59%	-5.95%	9.69%	\$ 52.29 MM		\$ 8.82 MM		560.000 MM	1,000,000,000
10	Golem	130	GNT	\$ 0.0311	0.00000432	-1.56%	-5.12%	-8.63%	-14.82%	-16.01%	-31.32%	-48.33%	-63.93%	\$ 30.45 MM		\$ 1.12 MM		980.050 MM	1,000,000,000
11	Qtum	41																96.29 MM	102,003,650

Figure 2: Fundamental analysis page of coincheckup.com

the day of the data snapshot. An assumption taken was that the market cap change would be of considerable importance. For cryptocurrencies, market cap is defined as the product of price and number of coins in circulation. However, it turns out the number of coins in circulation is modified or adjusted for a number of reasons. For instance, when it is discovered that some coins are lost forever (wallet seed is lost or destroyed) or the developers decide to “drop” privately held coins into circulation. As a result, this field was determined to be unreliable and removed from consideration for model building.

For the creation of a test and training set, the percentage change in price was calculated for each coin based on the December and February data. This value was plotted and based on the number of coins, a growth of greater than 30% was defined to be a successful outcome. A zoomed in plot of percentage change in price can be seen in Fig. 3. The red horizontal line marks the 30% cut-off of the price change. This was used to create a column in the data frame to set a binary value for said outcome.

Python files used to generate the dataset are attached in annex. We used *parse.py*, *parse-ia.py* and *parse-fa.py* to parse the HTML from 3 different pages of CoinCheckup.com, namely the fundamental analysis page (fa) and the investment analysis page (ia). It needs to be done for the version of the website in March and in December. Then we can merge all the data from March in *merge\_fixed\_data.py*.

Then in *add\_performance\_to\_merged\_data.py* we can read data from December and create 5 columns:

1. label\_Price: price in December
2. label\_growth\_rate\_Price:  $(\text{label\_Price} - \text{Price})/\text{Price}$
3. label\_Mkt. Cap: Market Cap in December

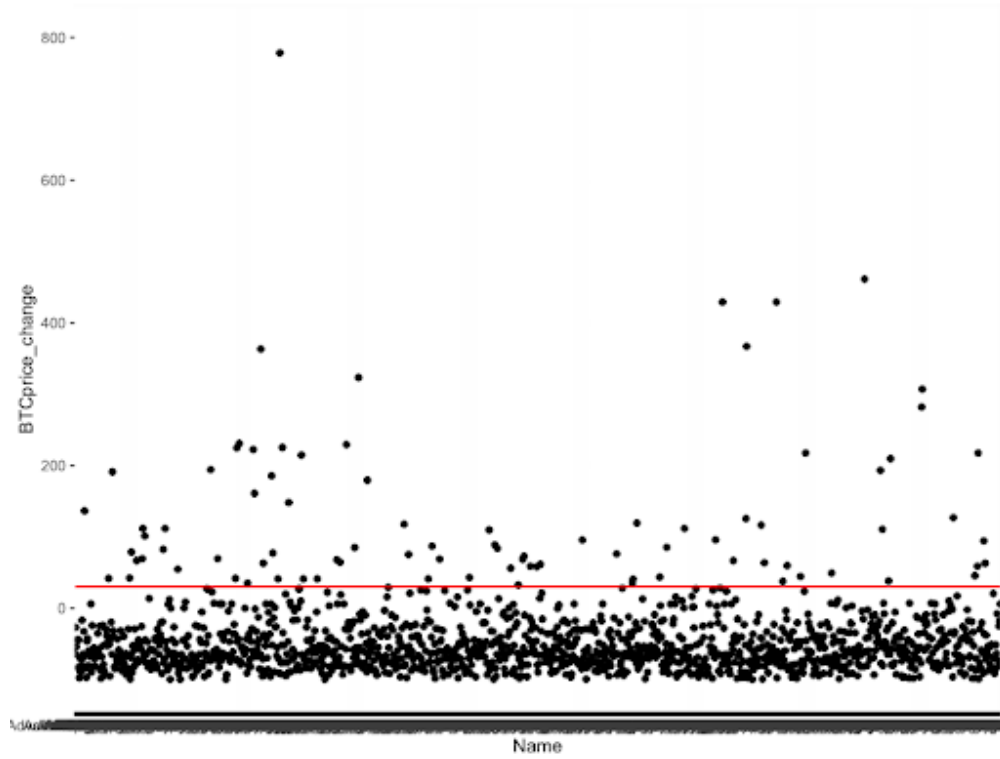


Figure 3: Plot of BTC Price Change for all coins

4.  $\text{label\_growth\_rate\_Mkt. Cap: } (\text{label\_Mkt. Cap} - \text{Mkt. Cap}) / \text{Mkt. Cap}$
5.  $\text{label\_disappeared}$ : whether a coin is listed in March but not in December

### 3 Dead coin prediction

#### 3.1 Logistic regression

The list of coins on Coincheckup.com is different between 10 months ago and now. Some have been added (new coins) and others have been removed. We label the removed coins as dead coins and the remaining coins as still alive. We want to predict the death of a coin over those 10 months based on the data that we had in March. As a result, we created a column "label\_disappeared" in the March dataset to tell whether the coin is dead or not. We split the dataset in 70% training and 30% testing (with fair distribution of dead coins in both parts).

#	Alive	Dead
Train	1025	208
Test	439	89

We begin with a logistic regression on the features which are numeric and have a very low rate of NAs. We check that our VIF is correct. Then we reduce the range of features to the relevant

ones and make a final logistic model, that we compare to the baseline We do a logistic regression with an arbitrary threshold  $p = 0.2$  and with an AIC of 1015.9 we get:

$$\hat{Y}_{dead} = \text{sigmoid}(-0.004X_{X45d} - 0.032X_{Social})$$

with  $\text{sigmoid} : x \rightarrow \frac{1}{1 + e^{-x}}$ . See appendix for feature descriptions.

We can make predictions on the test set and compare with the result 10 months later. The proportion of dead coins is around 20% and when we predict that if a coin is dead, it is true more or less half of the time, so the model is not too bad!

Real Pred	Alive	Dead	Model	Accuracy	TPR	FPR
Alive	320	117	Baseline	83.1%	0	0
Dead	37	52	Logistic reg.	70.5%	58.4%	26.8%

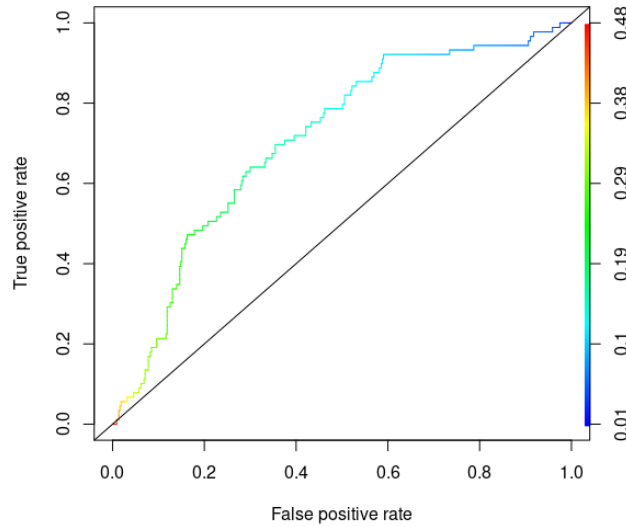


Figure 4: ROC of the logistic regression (dead coin prediction)

### 3.2 Random forest

We began checking the number of NAs in each column to select the features. We choose:  $X_{X24h}$ ,  $X_{X14d}$ ,  $X_{X45d}$ ,  $X_{Social}$ ,  $X_{Mkt.Cap}$ ,  $X_{Age.mo.}$  and  $X_{Business}$ . Then we try a random forest with cross-validation. We try a cross-validation measured with Accuracy and then F1-score as Accuracy is less relevant for unbalanced data.

Based on the results of figure 5, an mtry of 3 looks good. We can try sampling stratification to handle the imbalance of the data.

Real Pred	Alive	Dead	Model	Accuracy	TPR	FPR
Alive	427	10	Baseline	83.1%	0	0
Dead	70	19	Logistic reg.	70.5%	58.4%	26.8%
			Random forest	84.47%	21.35%	2.29%

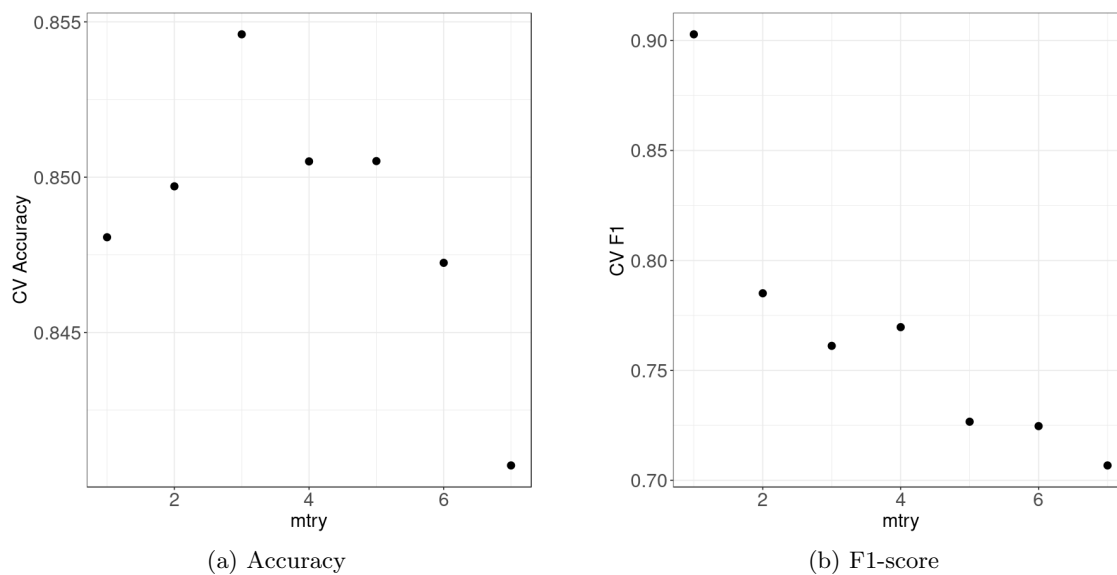


Figure 5: Cross validation of the random forest (dead coin prediction)

### 3.3 Gradient Boosting

Finally we can try a gradient boosting method with cross-validation.

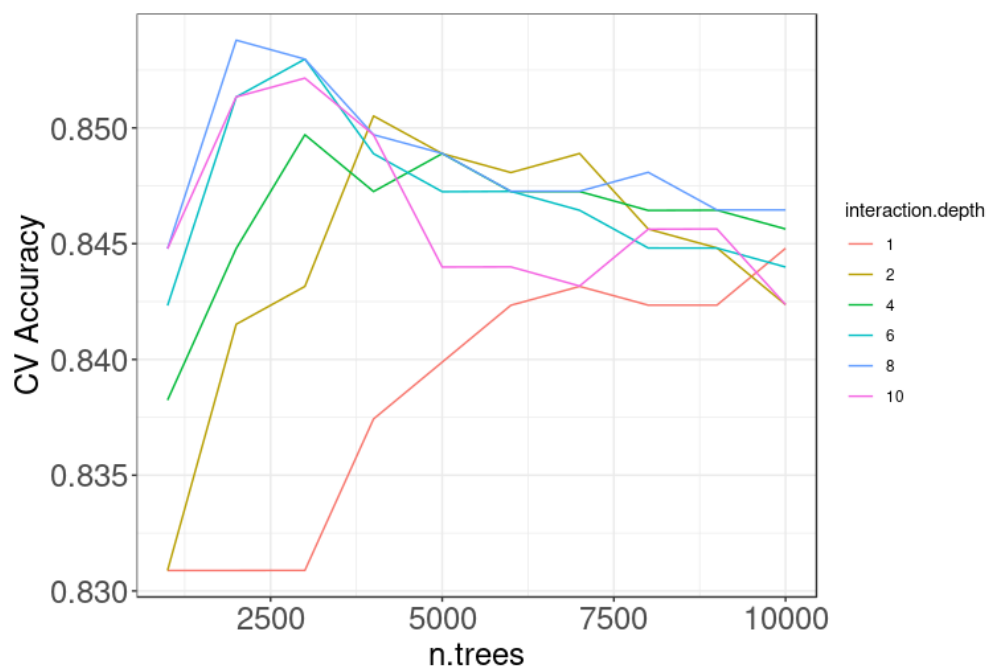


Figure 6: Cross validation of the gradient boosting (dead coin prediction)

Real Pred	Alive	Dead	Model	Accuracy	TPR	FPR
Alive	221	218	Baseline	83.1%	0	0
Dead	70	19	Logistic reg.	70.5%	58.4%	26.8%
			Random forest	84.47%	21.35%	2.29%
			Gradient boosting	45.45%	21.35%	49.66%

Based on results from figure 6, we choose `n.trees = 2000`, `interaction.depth = 8`.

### 3.4 Conclusion

Our random forest reaches the best accuracy, beating the baseline, while having a positive TPR and a low FPR.

Our gradient boosting model is very bad while the logistic regression threshold has been chosen to get the highest TPR, but it is at the expense of the accuracy and FPR.

## 4 Portfolio determination based on naive growth rates

Now let's tackle serious matters, we want to make money! To do that we need to establish a portfolio of assets. For now we assume that there is no opportunity cost for each asset to be in the portfolio. In practice there are fees to enter a position for each coin.

We want to naively predict  $Y = \text{label\_rate\_Price} = \frac{p_T - p_0}{p_0}$  over a period  $T$  of 10 months. After trying a linear regression on a lot of features, we stumble upon a multi-collinearity problem. So we remove the guilty features and then keep the few significant features:

$$Y = -2.53 \times 10^{-3} X_{Circ.Supply} + 9.58 \times 10^{-12} X_{Total.Supply} + \epsilon$$

Model	MSE	MAE	OSR <sup>2</sup>
Baseline	N/A	N/A	N/A
Linear reg.	547.6	1.94	-166.77

It looks bad. We can compare it to the baseline (predicting negative all the time). We assume that we invest the same amount of money in each coin (ignoring fees) if we follow the baseline while we invest sizes  $s_i = \frac{Y_i}{p_{0,i}}$  if we follow our regression strategy. We can thus compute the result growth rate of both portfolios:

$$1. \text{ Linear regression: } R = \frac{\sum_{i=1}^N s_i p_T^i - \sum_{i=1}^N s_i p_0^i}{\sum_{i=1}^N s_i p_0^i} = \frac{\sum_{i=1}^N Y_i r_i}{\sum_{i=1}^N Y_i} = 0.142$$

$$2. \text{ Baseline: } R_b = \frac{\sum_{i=1}^N r_i}{N} = -0.342$$

Surprisingly we still beat the baseline on the test set and even earn money. We are lucky.



## 5 Portfolio determination in a log normal stock model

### 5.1 Linear regression

We often assume that the price of a stock follows a log-normal distribution, that is to say that for  $S_t$  the price of an asset at time  $t$ ,  $\log(S_t) \sim \mathcal{N}(\mu, \sigma^2)$ . Let's  $B_t$  be a standard Brownian motion ( $\forall s < t, B_t - B_s \sim \mathcal{N}(0, t - s)$ ). We can write  $S_t = S_0 e^{\mu t + \sigma B_t}$ . Then for a time  $T > 0$  being our 10 months,  $Y^{log} = \log(S_T) - \log(S_0) = \mu T + \sigma B_T$  where  $B_T \sim \mathcal{N}(0, T)$ .

Thus performing a linear regression  $\forall 1 \leq i \leq n, Y_i^{log} = \beta_0 + \sum_{j=1}^p \beta_j X_{ij} + \epsilon_i$ , we can try, assuming our parameters constant, to estimate  $\mu = \frac{1}{T}(\beta_0 + \sum_{j=1}^p \beta_j X_{ij})$ . However estimating  $\sigma$  from  $\epsilon$  would require us to assume that each coin has the same constant volatility which is too strong.

We need to do a new split of the data. We come across a big multi-collinearity issue which forces us to remove some features. Making particular choices we find:

$$Y^{log} = \mu T = -1.06 - 0.00169X_{X7d} - 0.00033X_{MarketCap} + 0.00412X_{Age} + \epsilon$$

Model	MSE	MAE	OSR <sup>2</sup>	R
Baseline	1.35	0.774	0	0.413
Linear reg.	1.332	0.766	0.0152	0.50

We can create a portfolio with sizes  $s_i = \exp(Y^{log}_i)$ . Then our return is

$R_{log} = \frac{\sum_{i=1}^N s_i p_T^i - \sum_{i=1}^N s_i p_0^i}{\sum_{i=1}^N s_i p_0^i} = \frac{\sum_{i=1}^N s_i r_i}{\sum_{i=1}^N s_i} = 0.500$  against  $R_b = \frac{\sum_{i=1}^N r_i}{N} = 0.413$  for the baseline.

Much better already! We see that filtering out coins where information is missing, our portfolio behaves already better! That's because trustworthy coins get all their information filled.

### 5.2 Random forest

We take the features from the 2 previous models and we try it out on a random forest.

The best value is 1 for mtry.

Model	MSE	MAE	OSR <sup>2</sup>	R
Baseline	1.35	0.774	0	0.413
Linear reg.	1.332	0.766	0.0152	0.50
Random forest	1.343	0.768	0.005	0.318

Unfortunately this model does not beat the baseline in term of financial gain.

### 5.3 Gradient boosting

We use the same features as for the random forest. After cross-validation (figure 8, we conclude that the best parameters are the biggest ones.

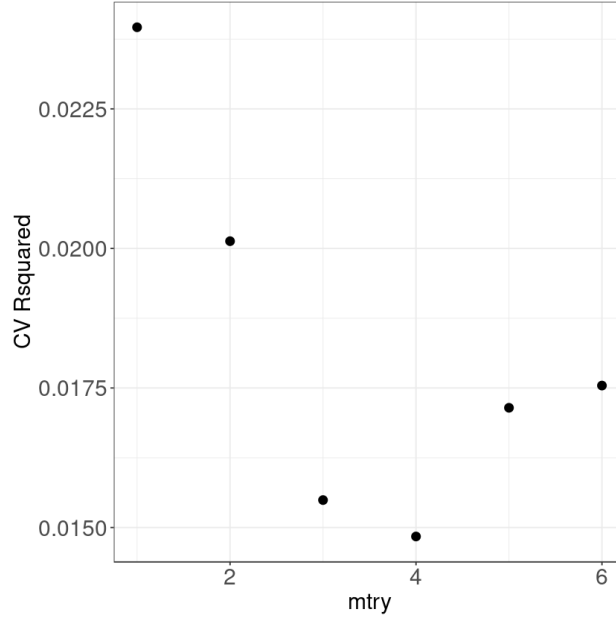


Figure 7: Cross validation of the random forest (log-normal model)

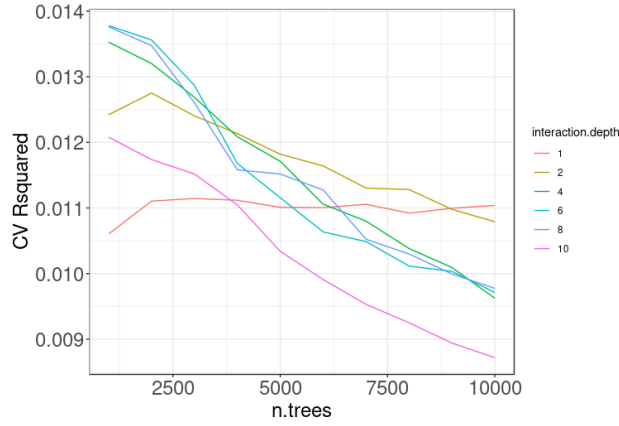


Figure 8: Cross validation of the gradient boosting (log-normal model)

Model	MSE	MAE	$OSR^2$	R
Baseline	1.35	0.774	0	0.413
Linear reg.	1.332	0.766	0.0152	0.50
Random forest	1.343	0.768	0.005	0.318
Gradient boosting	1.418	0.785	-0.050	0.290

This is also very bad compared to the baseline in term of portfolio return  $R$ .

## 5.4 Conclusion

The linear regression stays the best fit on the test set. The other model seem to overfit the data. Complicated models on low quality data makes bad results.

## 6 Portfolio determination in real-world

In reality there is an opportunity cost in choosing more coins in a portfolio. This cost comes from the fees to enter and exit the position: buy and sell fees plus withdraw fee (blockchain mining reward). Using the model from previous section. If we choose the coin with the best predicted  $\mu$  according to the linear regression, we get a growth rate  $r = 2.067$  for CannabisCoin, which outperforms  $R_{log}$ .

The top 5 includes: CannabisCoin, GlobalBoost-Y, Quebecoin, SolarCoin, Curecoin.

## 7 Appendix

### 7.1 Column Definitions

Column Name	Description
MC...past	Position of the coin the time of snapshot
Symbol_past	Trading symbol for each coin
Price_past	Price in USD at the time of snapshot
BTC_past	Price in BTC at the time of snapshot
X30d_past	Percentage change in price over 30 days
X45d_past	Percentage change in price over 45 days
X90d_past	Percentage change in price over 90 days
X200d_past	Percentage change in price over 200 days
Mkt..Cap_past	Market Cap in USD
MCAP.BTC_past	Market Cap in BTC
Circ..Supply_past	No of circulating coins in supply
Total.Supply_past	Total supply of coins
Max..Supply_past	Maximum supply of coins possible
Team_past	Confidence value of the team based on past performance of leadership, developers etc, in percentage.
Advisors_past	Confidence value of the advisors, whether they have been successful historically and confirmed to be associated with said coin, in percentage.
Brand.Buzz_past	How large and active is the community as compared to the rest of the market, in percentage
Product_past	Is the coin an idea or a product at this stage, does it have a roadmap and available, confidence in percentage
Coin_past	Confidence in ability to transact high volumes and availability, in percentage
Social_past	Confidence in activity over social media, in percentage
Communication_past	Confidence of accessibility and activity over slack, telegram, email etc, in percentage
Business_past	Confidence in investors behind project, publishing of revenue reports etc, in percentage
GitHub_past	Confidence in activity on Github, in percentage
Age..mo._past	Age of the coins, in number of months
Winning.months_past	No of successful months with positive growth over 12 months, in percentage
Start.price_past	ICO start price
CMGR....3mo_past	Compound Monthly Growth Rate trailing 3 months
Cum..ROI_past	Cumulative ROI, the ROI in % from start trading price till the current price
Avg..volume_past	Average volume of coins traded
ATH_past	All time high of coin in USD

X..fm..ATH_past	Change in USD price from ATH, in percentage
ATH..BTC._past	All time high of coin in BTC
X..fm..ATH..BTC._past	Change in BTC price from ATH, in percentage
Name	Name of the coin
Mkt.BTC_c	Market cap conversion from USD to BTC
growth8mth	Price change over month duration
BTCprice_change	Percentage change in BTC compared to December
BTCmcap_change	Market cap change of in BTC compared to December snapshot
success	Binary value where 1 equals greater than 30% increase in price

## 7.2 Python Code

*parse.py*

---

```
# URL = "https://web.archive.org/web/20190208144932/https://coincheckup.com/predictions"

import re
from bs4 import BeautifulSoup
import csv

datas = []
for tab_idx in range(1, 22):
    with open(f"html/20190208/predictions/Predictions Overview –
CoinCheckup{tab_idx}.html") as file:
        soup = BeautifulSoup(file.read(), 'html.parser')

        cryptos = [link.string for link in soup.find_all('a') if link.get("href") and "coins"
in link.get("href")]

        divs = [row for row in soup.find_all(attrs={"class": re.compile("ag-row
ag-row-no-focus ag-row-\w* ag-row-level-0")}) if
row.div is not None and row.div.span is None]
        data = [[cryptos[idx]] + [cell.string.strip("$").replace(',', '') for cell in
div.children] for idx, div in enumerate(divs)]

        datas.append(data)

with open('coincheckup_10_months_predictions.csv', 'w', newline='') as csvfile :
    writer = csv.writer( csvfile , delimiter=',')
    writer.writerow([span.string for span in soup.find_all (role="columnheader")])
    for data in datas:
        for row in data:
            writer.writerow(row)
```

---

---

```

# URL = "https://web.archive.org/web/20190208144932/https://coincheckup.com/analysis"

import re
from bs4 import BeautifulSoup
import csv

def current():
    datas = []
    for tab_idx in range(1, 21):
        with open(f"html/20191205/fundamental_analysis/Analysis Overview -
CoinCheckup{tab_idx}.html") as file:
            soup = BeautifulSoup(file.read(), 'html.parser')

            cryptos = [link.string for link in soup.find_all('a') if link.get("href") and
"coins" in link.get("href")]

            divs = [row for row in soup.find_all(attrs={"class": re.compile("ag-row
ag-row-no-focus ag-row-\w* ag-row-level-0")}) if
row.div is not None and row.contents[1].img is None]
            data = [[cryptos[idx]] + [div.contents[i].string if div.contents[i].div is None
else div.contents[i].div.string for i in range(len(div.contents)-1) if i != 1]
for idx, div in enumerate(divs)]
            for d in data:
                #print(d)
                for i in range(len(d)):
                    if d[i] is not None: # Some 200d values are None
                        d[i] = d[i].strip("$% ") # MM and K units should be handled here
            datas.append(data)

with open('coincheckup_current_fa.csv', 'w', newline='') as csvfile :
    writer = csv.writer( csvfile , delimiter=',')
    ### WARNING: Problem with order in the column, FIXED HEADER
    header=['Name', 'MC #', 'Symbol', 'Price', 'BTC', '1h', '24h', '7d', '14d', '30d',
'45d', '90d', '200d', 'Mkt. Cap', 'MCAP BTC', '24h Vol', '24h Vol BTC', 'Circ.
Supply', 'Total Supply', 'Max. Supply', 'Team', 'Advisors', 'Brand/Buzz',
'Product', 'Coin', 'Social', 'Communication', 'Business', 'GitHub', 'GitHub', 'Avg.
volume', 'Age (mo)', 'Winning months']
    #writer.writerow([span.string for span in soup.find_all(role="columnheader")[1:-1] if
span.string is not None])
    writer.writerow(header)
    for data in datas:
        for row in data:

```

```

writer.writerow(row)

def ten_months():
    datas = []
    for tab_idx in range(1, 19):
        with open(f"html/20190208/fundamental_analysis/Analysis Overview –
CoinCheckup{tab_idx}.html") as file:
            soup = BeautifulSoup(file.read(), 'html.parser')

            cryptos = [link.string for link in soup.find_all('a') if link.get("href") and
"coins" in link.get("href")]

            divs = [row for row in
                    soup.find_all(attrs={"class": re.compile("ag-row ag-row-no-focus
ag-row-\w* ag-row-level-0")}) if
                    row.div is not None and row.contents[1].img is None]
            data = [
                [cryptos[idx]] + [div.contents[i].string if div.contents[i].div is None else
div.contents[i].div.string
                    for i in range(len(div.contents) - 2)] for idx, div in
                    enumerate(divs)]

            for d in data:
                print(d)
                for i in range(len(d)):
                    ### WARNING: I modify this part because there is a column that is
                    useless
                    if i <=2 :
                        if i == 2 :
                            continue
                        if d[i] is not None: # Some 200d values are None
                            d[i] = d[i].strip("% ")

                    else :
                        if d[i] is not None: # Some 200d values are None
                            d[i-1] = d[i].strip("% ")

            datas.append(data)

with open('coincheckup_10_months_fa.csv', 'w', newline='') as csvfile :
    writer = csv.writer( csvfile , delimiter=',')

    ### WARNING: Problem with order in the column, FIXED HEADER
    header=['Name', 'MC #', 'Symbol', 'Price', 'BTC', '1h', '24h', '7d', '14d', '30d',

```



```

'45d', '90d', '200d', 'Mkt. Cap', 'MCAP BTC', '24h Vol', '24h Vol BTC', 'Circ.
Supply', 'Total Supply', 'Max. Supply', 'Team', 'Advisors', 'Brand/Buzz',
'Product', 'Coin', 'Social', 'Communication', 'Business', 'GitHub', 'GitHub', 'Avg.
volume', 'Age (mo)', 'Winning months']
#writer.writerow([span.string for span in soup.find_all (role="columnheader") [1:-1] if
span.string is not None])
writer.writerow(header)
for data in datas:
    for row in data:
        writer.writerow(row)

current()
ten_months()

```

---

```
# URL = "https://web.archive.org/web/20190208144932/https://coincheckup.com/analysis"
```

```
import re
from bs4 import BeautifulSoup
import csv
```

```
def current():
    datas = []
    numeration=1
    for tab_idx in range(1, 25):
        with open(f'html/20191205/investment_analysis/Investment Overview –
CoinCheckup{tab_idx}.html') as file:
            soup = BeautifulSoup(file.read(), 'html.parser')

            cryptos = [link.string for link in soup.find_all('a') if link.get("href") and
"coins" in link.get("href")]

            divs = [row for row in soup.find_all(attrs={"class": re.compile("ag-row
ag-row-no-focus ag-row-\w* ag-row-level-0")}) if
row.div is not None and row.contents[1].img is None]

            data = [[cryptos[idx]] + [ cell.string.strip ("% % ").replace(', ', ' ') for cell in
div.children ] for idx, div in enumerate(divs)]

            datas.append(data)
```

```
with open('coincheckup_current_ia.csv', 'w', newline='') as csvfile :
    writer = csv.writer( csvfile , delimiter=',')
    writer.writerow([span.string for span in soup.find_all (role="columnheader")][1:] if
span.string is not None])
    for data in datas:
        for row in data:
            data.insert
            writer.writerow(row)
```

```
def ten_months():
    datas = []
    for tab_idx in range(1, 22):
        with open(f'html/20190208/investment_analysis/Investment Overview –
CoinCheckup{tab_idx}.html') as file:
```

```

soup = BeautifulSoup(file.read(), 'html.parser')

cryptos = [link.string for link in soup.find_all('a') if link.get("href") and
           "coins" in link.get("href")]

divs = [row for row in
        soup.find_all(attrs={"class": re.compile("ag-row ag-row-no-focus
        ag-row-\w* ag-row-level-0")}) if
        row.div is not None and row.contents[1].img is None]
if len(cryptos) == 0:
    print("WARNING: File of idx {} is ignored because it contains no
    data".format(tab_idx))
    continue
data = [[cryptos[idx]] + [ cell.string.strip ("%%" ).replace(',',' ') for cell in
        div.children] for idx, div in enumerate(divs)]

datas.append(data)

with open('coincheckup_10_months_ia.csv', 'w', newline='') as csvfile :
    writer = csv.writer( csvfile , delimiter=',')
    writer.writerow([span.string for span in soup.find_all (role="columnheader")[1:] if
        span.string is not None])
    for data in datas:
        for row in data:
            writer.writerow(row)

current()
ten_months()

```

---

```
import csv

num_fields = ['MC #', 'Price', 'BTC', '1h', '24h', '7d', '14d', '30d', '45d', '90d', '200d',
              'Mkt. Cap', 'MCAP BTC', '24h Vol', '24h Vol BTC', 'Circ. Supply', 'Total Supply', 'Max.
              Supply', 'Team', 'Advisors', 'Brand/Buzz', 'Product', 'Coin', 'Social', 'Communication',
              'Business', 'GitHub', 'Avg. volume', 'Age (mo)', 'Winning months']

with open("fixed_data/coincheckup_10_months_fa.csv", 'r', newline='') as old_fa,
    open("fixed_data/coincheckup_10_months_fa.csv", 'r', newline='') as old_ia:
    fa_reader = csv.DictReader(old_fa, delimiter=',')
    ia_reader = csv.DictReader(old_ia, delimiter=',')

    fields = set(fa_reader.fieldnames).union(set(ia_reader.fieldnames))
    #data = pd.DataFrame({field: [] for field in fields})
    data = {row["Name"]: row for row in fa_reader}
    for row in ia_reader:
        for key, value in row.items():
            data[row["Name"]][key] = value

    for name, row in data.items():
        for field in num_fields:
            if row[field] in {'--', '---', '', 'N/A'}:
                row[field] = ''
            else:
                row[field] = float(row[field].strip("$% ").replace(',', ''),
                                   '').replace('K', '000').replace('MM', '000000').replace('Bn',
                                   '000000000').replace('Tn', '000000000000'))

with open("merged_10months_data.csv", 'w', newline='') as output:
    writer = csv.DictWriter(output, delimiter=',', fieldnames=fields)
    writer.writeheader()
    for row in data.values():
        writer.writerow(row)
```

---

## 7.3 Annotated R Code

### 7.3.1 Dead coin prediction

```
In [2]: library(caret)
library(MASS)
library(caTools)
library(randomForest)
library(ggplot2)
library(GGally)
library(car)
library(rpart)
library(rattle)
library(boot)
library(dplyr)
library(ROCR)

mean_squared_error <- function(responses, predictions) {
  MSE <- mean(((responses - predictions))^2)
  return(MSE)
}

mean_absolute_error <- function(responses, predictions) {
  MAE <- mean(abs(responses - predictions))
  return(MAE)
}

OS_R_squared <- function(responses, predictions, train_responses) {
  baseline <- mean(train_responses)
  SSE <- sum((responses - predictions)^2)
  SST <- sum((responses - baseline)^2)
  r2 <- 1 - SSE/SST
  return(r2)
}

all_metrics <- function(responses, predictions, train_responses) {
  filter_vec = !is.na(responses) & !is.na(predictions)
  responses <- responses[filter_vec]
  predictions <- predictions[filter_vec]
  train_responses <- train_responses[filter_vec]
  mse <- mean_squared_error(responses, predictions)
  mae <- mean_absolute_error(responses, predictions)
  OSR2 <- OS_R_squared(responses, predictions, train_responses)
  return(c(mse, mae, OSR2))
}

tableAccuracy <- function(label, pred) {
  t = table(label, pred)
```

```

    a = sum(diag(t))/length(label)
    return(a)
}

tableTPR <- function(label, pred) {
  t = table(label, pred)
  return(t[2,2]/(t[2,1] + t[2,2]))
}

tableFPR <- function(label, pred) {
  t = table(label, pred)
  return(t[1,2]/(t[1,1] + t[1,2]))
}

```

We load the preprocessed data:

```

In [3]: df <- read.csv("../python/coincheckup/dataset.csv", stringsAsFactors=FALSE, na="")
df$label_disappeared = as.factor(df$label_disappeared)
# name and Symbol are the only strings

```

#### a) Logistic regression

We begin with a logistic regression on the features which are numeric and have a very low rate of NAs:

```

In [4]: set.seed(42)
split = sample.split(df$label_disappeared, SplitRatio = 0.7)
train_data <- filter(df, split== TRUE)
test_data <- filter(df, split== FALSE)

table(train_data$label_disappeared)
table(test_data$label_disappeared)

#LOGISTIC
logistic <- glm(label_disappeared ~ Price + X1h + X24h + X7d + X14d + X30d + X45d +
  X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
  Team + Product + Coin + Social + Communication + Business +
  Avg..volume + Age..mo., data = train_data, family="binomial")
summary(logistic)

pred = predict(logistic, newdata = test_data, type="response")

table(test_data$label_disappeared, pred>0.2)

tableAccuracy(test_data$label_disappeared, pred>0.2)
tableTPR(test_data$label_disappeared, pred>0.2)
tableFPR(test_data$label_disappeared, pred>0.2)

vif(logistic)

```

```
FALSE TRUE
1025 208
```

```
FALSE TRUE
439 89
```

Warning message:

“glm.fit: fitted probabilities numerically 0 or 1 occurred”

Call:

```
glm(formula = label_disappeared ~ Price + X1h + X24h + X7d +
    X14d + X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol +
    Circ..Supply + Total.Supply + Team + Product + Coin + Social +
    Communication + Business + Avg..volume + Age..mo., family = "binomial",
    data = train_data)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.4129	-0.6457	-0.4438	-0.2806	2.6774

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	4.904e-02	4.592e-01	0.107	0.91494
Price	-5.873e-03	1.089e-02	-0.539	0.58957
X1h	7.936e-05	5.658e-03	0.014	0.98881
X24h	4.684e-04	2.211e-03	0.212	0.83221
X7d	2.764e-03	1.811e-03	1.526	0.12690
X14d	-5.093e-04	1.361e-03	-0.374	0.70821
X30d	5.891e-04	2.114e-03	0.279	0.78049
X45d	-6.051e-03	2.321e-03	-2.608	0.00912 **
X90d	2.879e-03	1.871e-03	1.539	0.12384
X200d	-2.364e-03	2.256e-03	-1.048	0.29464
Mkt..Cap	-5.270e-05	3.726e-04	-0.141	0.88753
X24h.Vol	-4.836e-04	4.349e-04	-1.112	0.26623
Circ..Supply	5.619e-05	4.204e-04	0.134	0.89366
Total.Supply	2.683e-14	4.072e-14	0.659	0.51002
Team	-2.397e-02	1.576e-02	-1.520	0.12847
Product	-4.131e-03	4.869e-03	-0.848	0.39625
Coin	7.949e-03	5.937e-03	1.339	0.18066
Social	-1.858e-02	8.662e-03	-2.145	0.03199 *
Communication	-3.273e-03	3.600e-03	-0.909	0.36335
Business	-3.909e-04	3.243e-03	-0.121	0.90407
Avg..volume	-8.158e-11	8.827e-10	-0.092	0.92636

```
Age..mo.      7.262e-03  5.277e-03  1.376  0.16878
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

(Dispersion parameter for binomial family taken to be 1)

```
Null deviance: 1087.97 on 1191 degrees of freedom
Residual deviance: 954.15 on 1170 degrees of freedom
(41 observations deleted due to missingness)
AIC: 998.15
```

Number of Fisher Scoring iterations: 11

	FALSE	TRUE
FALSE	307	117
TRUE	31	55

```
0.685606060606061
0.63953488372093
0.275943396226415
Price 1.02584781560209 X1h 1.15769836383294 X24h 2.96033449373743 X7d 2.95185035557887
X14d 3.24090275717362 X30d 4.91817668782566 X45d 4.59783289757329 X90d 2.61198284261179
X200d 1.91331591388318 Mkt..Cap 1.03647241353168 X24h.Vol 1.03491142934378 Circ..Supply
1.05018223690908 Total.Supply 1.03860840602233 Team 7.10648538827917 Product
1.81658646390026 Coin 2.19831931567567 Social 3.19603310898613 Communication
1.46875216267816 Business 1.5922862721872 Avg..volume 1.0164595410715 Age..mo.
1.3693268345375
```

Then we reduce the range of features to the relevant ones and make a final logistic model, that we compare to the baseline:

```
In [5]: logistic <- glm(label_disappeared ~ X45d + Social + 0, data = train_data,
      family="binomial")
summary(logistic)

pred = predict(logistic, newdata = test_data, type="response")

table(test_data$label_disappeared, pred>0.2)

tableAccuracy(test_data$label_disappeared, pred>0.2)
tableTPR(test_data$label_disappeared, pred>0.2)
tableFPR(test_data$label_disappeared, pred>0.2)

# Baseline accuracy:
t_ <- table(test_data$label_disappeared)
```



```

t_[1]/sum(t_)

rocr.pred <- prediction(pred, test_data$label_disappeared)
perf <- performance(rocr.pred, "tpr", "fpr")
plot(perf, colorize = TRUE)
abline(0, 1)
as.numeric(performance(rocr.pred, "auc")@y.values)

```

Warning message:

“glm.fit: fitted probabilities numerically 0 or 1 occurred”

Call:

```

glm(formula = label_disappeared ~ X45d + Social + 0, family = "binomial",
    data = train_data)

```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-1.1773	-0.6347	-0.4854	-0.3496	2.7011

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
X45d	-0.004062	0.001108	-3.667	0.000245 ***
Social	-0.032039	0.001607	-19.935	< 2e-16 ***

---

Signif. codes: 0 ‘\*\*\*’ 0.001 ‘\*\*’ 0.01 ‘\*’ 0.05 ‘.’ 0.1 ‘ ’ 1

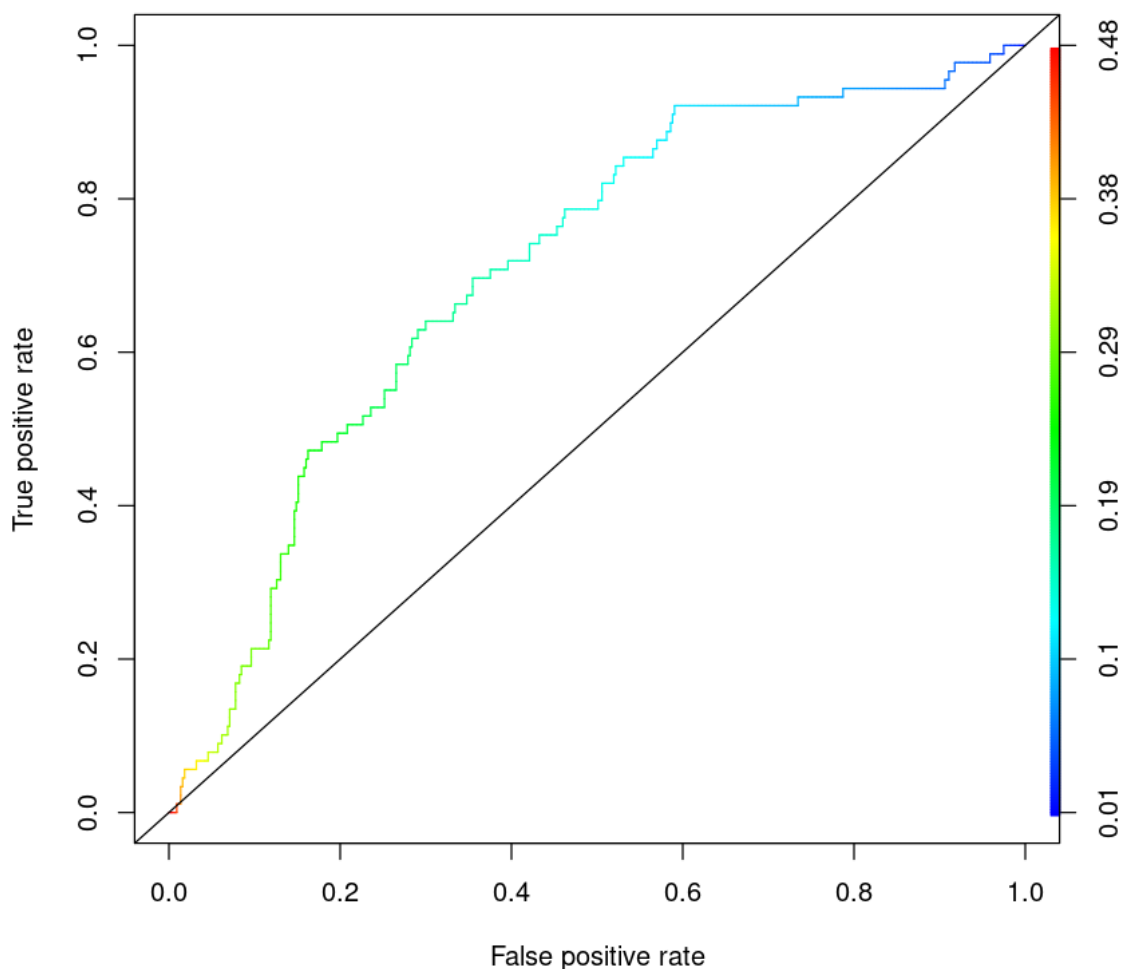
(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1696.8 on 1224 degrees of freedom  
Residual deviance: 1011.9 on 1222 degrees of freedom  
(9 observations deleted due to missingness)  
AIC: 1015.9

Number of Fisher Scoring iterations: 8

	FALSE	TRUE
FALSE	320	117
TRUE	37	52

0.704545454545455  
0.584269662921348  
0.267734553775744  
**FALSE:** 0.831439393939394  
0.708919342812334



## b) Random Forest

Let's begin checking the number of NAs in each column to select the features.

```
In [10]: sapply(df, function(x){ sum(is.na(x)) })
```

```
Name 0 MC.. 0 Symbol 0 Price 0 BTC 0 X1h 0 X24h 0 X7d 0 X14d 0 X30d 4 X45d 11 X90d 22
X200d 48 Mkt..Cap 0 MCAP.BTC 0 X24h.Vol 0 X24h.Vol.BTC 0 Circ..Supply 0 Total.Supply 9
Max..Supply 1394 Team 0 Advisors 566 Brand.Buzz 566 Product 0 Coin 0 Social 0
Communication 0 Business 0 GitHub 697 GitHub.1 697 Avg..volume 0 Age..mo. 1
Winning.months 1 label_Price 297 label_Mkt..Cap 297 label_growth_rate_Price 297
label_growth_rate_Mkt..Cap 297 label_disappeared 0
```

Now let's try a random forest with cross-validation. The dataset is not balanced so we try a cross-validation measured with Accuracy and then F1-score as Accuracy is less relevant in this case.

```
In [11]: set.seed(42)
test_data_filled_with_0 <- test_data
test_data_filled_with_0[is.na(test_data_filled_with_0)] <- 0
train_data.mm = as.data.frame(model.matrix(label_disappeared ~ X24h + X14d + X45d +
  Social + Mkt..Cap + Age..mo. + Business, data = train_data))
test_data.mm = as.data.frame(model.matrix(label_disappeared ~ X24h + X14d + X45d +
  Social + Mkt..Cap + Age..mo. + Business, data = test_data_filled_with_0))

train.rf <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business,
  data = train_data,
  method = "rf",
  na.action = na.omit,
  tuneGrid = data.frame(mtry=1:7),
  trControl = trainControl(method="cv", number=5),
  metric = "Accuracy")

train.rf$results
train.rf
best.rf <- train.rf$finalModel
pred.rf <- predict(best.rf, newdata = test_data.mm) # can use same model matrix

ggplot(train.rf$results, aes(x = mtry, y = Accuracy)) + geom_point(size = 3) +
  ylab("CV Accuracy") + theme_bw() + theme(axis.title=element_text(size=18),
    axis.text=element_text(size=18))
```

mtry <int>	Accuracy <dbl>	Kappa <dbl>	AccuracySD <dbl>	KappaSD <dbl>
1	0.8480651	0.1820981	0.01230487	0.07219353
2	0.8497045	0.2675792	0.01730529	0.08837830
3	0.8545991	0.3131013	0.01457905	0.07484737
4	0.8505074	0.2927120	0.01216925	0.05062367
5	0.8505175	0.3181606	0.01943042	0.07520890
6	0.8472421	0.2995333	0.01863631	0.06998712
7	0.8407181	0.2733978	0.02110235	0.08098173

Random Forest

```
1233 samples
  7 predictor
  2 classes: 'FALSE', 'TRUE'
```

No pre-processing

Resampling: Cross-Validated (5 fold)

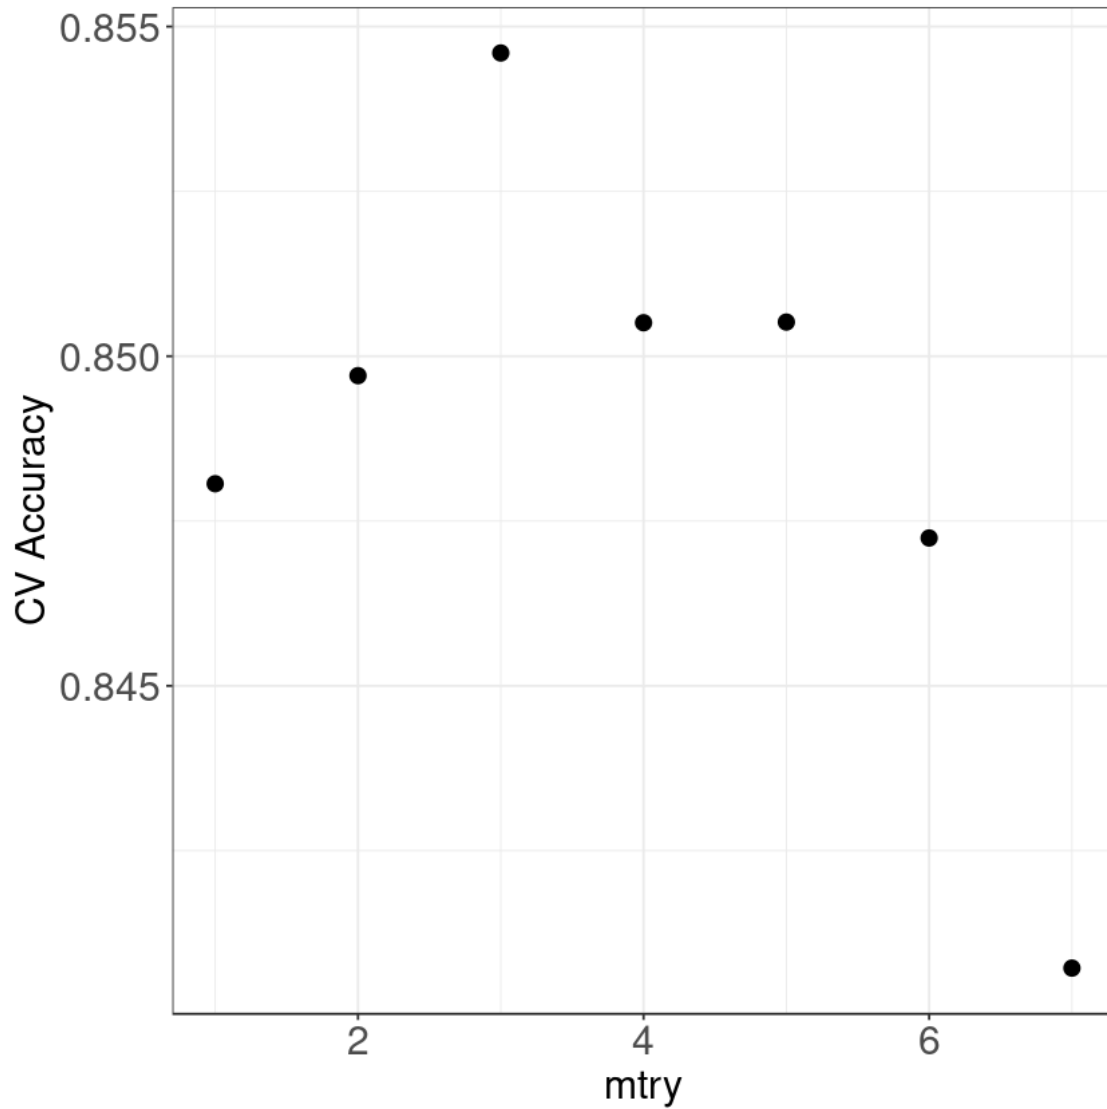
Summary of sample sizes: 980, 978, 980, 979, 979

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
1	0.8480651	0.1820981
2	0.8497045	0.2675792

3	0.8545991	0.3131013
4	0.8505074	0.2927120
5	0.8505175	0.3181606
6	0.8472421	0.2995333
7	0.8407181	0.2733978

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was `mtry = 3`.



```
In [14]: f1 <- function (data, lev = NULL, model = NULL) {
  pred <- data$pred[!is.na(data$pred)&!is.na(data$obs)]
  obs <- data$obs[!is.na(data$pred)&!is.na(data$obs)]
```

```

precision <- posPredValue(pred, obs, positive = lev[2])
recall <- sensitivity(pred, obs, postive = lev[1])
f1_val <- (2 * precision * recall) / (precision + recall)
names(f1_val) <- c("F1")
#print(precision)
#print(recall)
#print(f1_val)
f1_val
}

train.rf <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business,
  data = train_data,
  method = "rf",
  na.action = na.omit,
  tuneGrid = data.frame(mtry=1:7),
  trControl = trainControl(method="cv", number=5, summaryFunction = f1)
  metric = PY1+s"F1")

train.rf$results
train.rf
best.rf <- train.rf$finalModel
pred.rf <- predict(best.rf, newdata = test_data.mm) # can use same model matrix

ggplot(train.rf$results, aes(x = mtry, y = F1)) + geom_point(size = 3) +
  ylab("CV F1") + theme_bw() + theme(axis.title=element_text(size=18),
  axis.text=element_text(size=18))

```

mtry <int>	F1 <dbl>	F1SD <dbl>
1	0.9028305	0.08709772
2	0.7850726	0.07805387
3	0.7611599	0.11121317
4	0.7696718	0.06996851
5	0.7266709	0.10382716
6	0.7246719	0.11886381
7	0.7067876	0.12034424

Random Forest

1233 samples

7 predictor

2 classes: 'FALSE', 'TRUE'

No pre-processing

Resampling: Cross-Validated (5 fold)

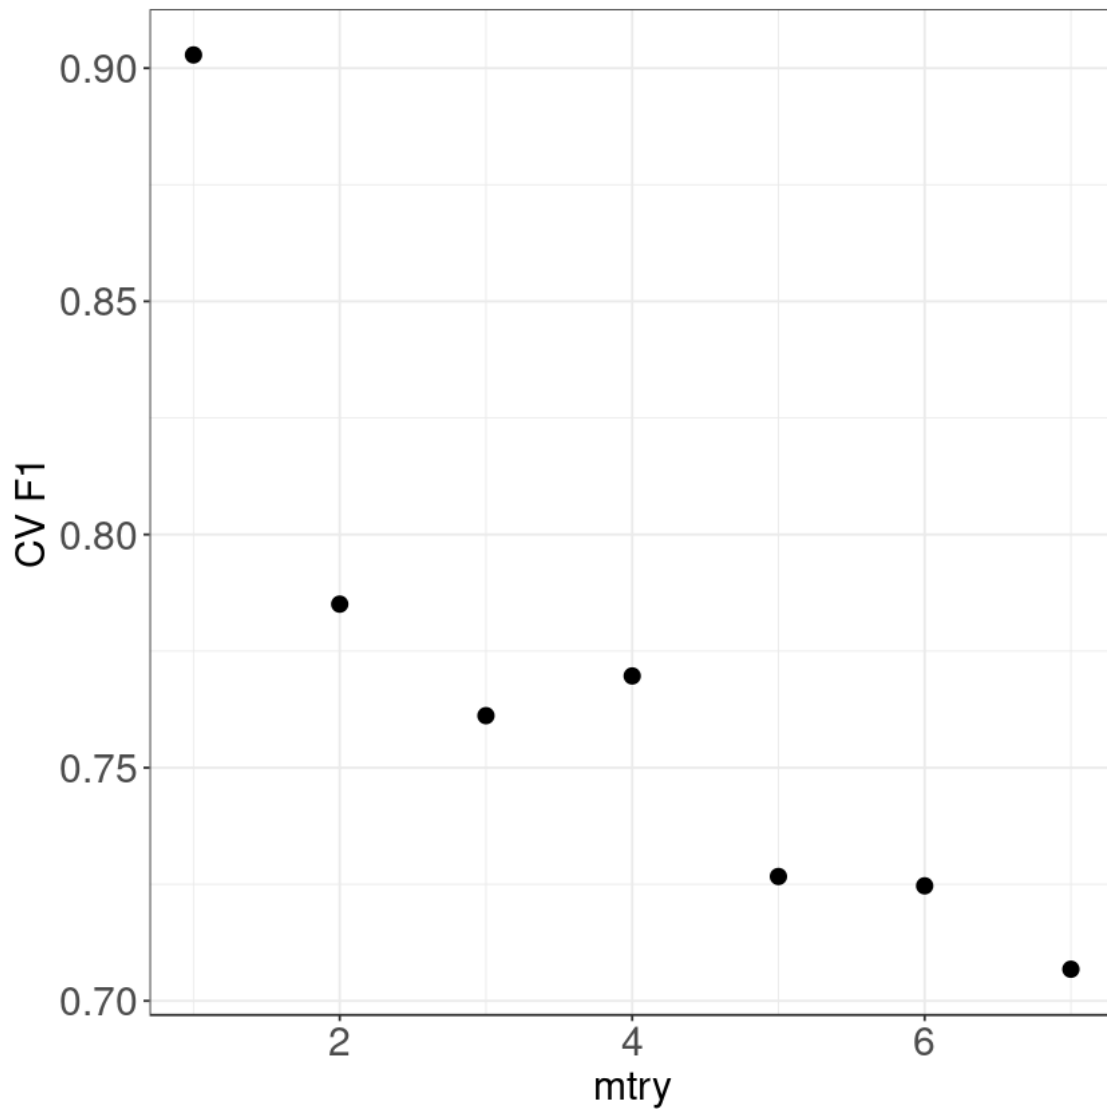
Summary of sample sizes: 979, 980, 979, 978, 980

Resampling results across tuning parameters:

mtry F1

1	0.9028305
2	0.7850726
3	0.7611599
4	0.7696718
5	0.7266709
6	0.7246719
7	0.7067876

F1 was used to select the optimal model using the largest value.  
The final value used for the model was `mtry = 1`.



An `mtry` of 3 looks good. We can try sampling stratification to handle the imbalance of the data.

```
In [15]: set.seed(42)
rare.class.prevalence = 0.2
nRareSamples = 1000 * rare.class.prevalence
mod.rf <- randomForest(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business, data = train_data, mtry = 3, nodesize = 5, ntree = 1000,
  strata=train_data$label_disappeared,
  sampsize=c(nRareSamples,nRareSamples), na.action = na.omit)
# print(mod.rf)

pred.rf <- predict(mod.rf, newdata = test_data)

table(test_data$label_disappeared, pred.rf)

tableAccuracy(test_data$label_disappeared, pred.rf)
tableTPR(test_data$label_disappeared, pred.rf)
tableFPR(test_data$label_disappeared, pred.rf)

pred.rf
  FALSE TRUE
FALSE  427  10
TRUE   70  19

0.84469696969697
0.213483146067416
0.022883295194508
```

### c) Boosting

Finally we can try a boosting method with cross-validation.

```
In [17]: tGrid = expand.grid(n.trees = (1:10)*1000, interaction.depth = c(1,2,4,6,8,10),
  shrinkage = 0.001, n.minobsinnode = 10)

set.seed(42)
train.boost <- train(label_disappeared ~ X24h + X14d + X45d + Social + Mkt..Cap +
  Age..mo. + Business,
  data = train_data,
  method = "gbm", ## gradient boosting machine
  tuneGrid = tGrid,
  trControl = trainControl(method="cv", number=5),
  metric = "Accuracy",
  na.action = na.omit)

train.boost
best.boost <- train.boost$finalModel

ggplot(train.boost$results, aes(x = n.trees, y = Accuracy, colour =
  as.factor(interaction.depth))) + geom_line() +
  ylab("CV Accuracy") + theme_bw() + theme(axis.title=element_text(size=18),
```

```
axis.text=element_text(size=18)) +  
scale_color_discrete(name = "interaction.depth")
```

Stochastic Gradient Boosting

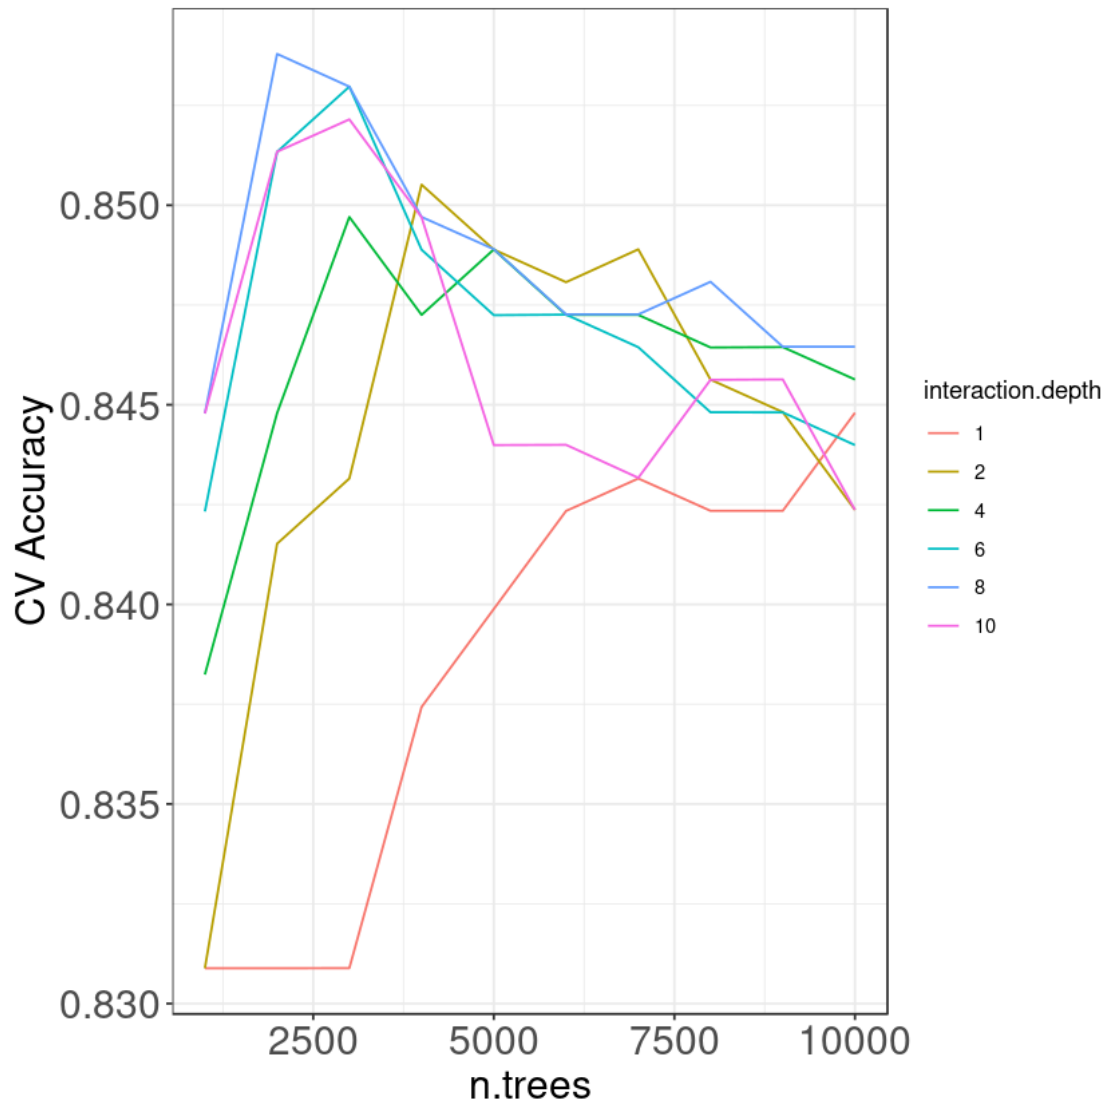
1233 samples  
7 predictor  
2 classes: 'FALSE', 'TRUE'

No pre-processing  
Resampling: Cross-Validated (5 fold)  
Summary of sample sizes: 980, 978, 980, 979, 979

Tuning parameter 'shrinkage' was held constant at a value of 0.001

Tuning parameter 'n.minobsinnode' was held constant at a value of 10  
Accuracy was used to select the optimal model using the largest value.  
The final values used for the model were n.trees = 2000, interaction.depth =  
8, shrinkage = 0.001 and n.minobsinnode = 10.





We choose `n.trees = 2000`, `interaction.depth = 8`.

```
In [18]: pred.best.boost <- predict(best.boost, newdata = test_data.mm, n.trees = 2000,
  interaction.depth = 8) # from CV plot
```

```
table(test_data$label_disappeared, pred.best.boost>0.7)
```

```
tableAccuracy(test_data$label_disappeared, pred.best.boost>0.7)
```

```
tableTPR(test_data$label_disappeared, pred.best.boost>0.7)
```

```
tableFPR(test_data$label_disappeared, pred.best.boost>0.7)
```

```
FALSE TRUE
FALSE  221  218
```

TRUE      70    19

0.454545454545455  
0.213483146067416  
0.496583143507973

### 7.3.2 Portfolio determination based on naive growth rates

#### a) Regression

We try a linear model on a lot of features.

```
In [22]: set.seed(42)
split = sample.split(df$label_growth_rate_Price, SplitRatio = 0.7)
price_train_data <- filter(df, split== TRUE)
price_test_data <- filter(df, split== FALSE)

price_train_data = price_train_data %>% filter(as.integer(label_disappeared)==1)
# we keep label_disappeared = FALSE
price_test_data = price_test_data %>% filter(as.integer(label_disappeared)==1)

mod <- lm(label_growth_rate_Price ~ Price + X1h + X24h + X7d + X14d + X30d + X45d +
          X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
          Team + Product + Coin + Social + Communication + Business +
          Avg..volume + Age..mo., data = price_train_data)

summary(mod)

pred_lm = predict(mod, newdata = price_test_data)

all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
            price_train_data$label_growth_rate_Price)
```

Call:

```
lm(formula = label_growth_rate_Price ~ Price + X1h + X24h + X7d +
    X14d + X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol +
    Circ..Supply + Total.Supply + Team + Product + Coin + Social +
    Communication + Business + Avg..volume + Age..mo., data = price_train_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-95.323	-0.448	0.006	0.504	157.521

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-1.380e+00	1.463e+00	-0.943	0.346
Price	-1.363e-05	9.791e-04	-0.014	0.989

X1h	2.563e-02	2.873e-02	0.892	0.373
X24h	-1.096e-02	9.764e-03	-1.122	0.262
X7d	-3.442e-03	7.134e-03	-0.482	0.630
X14d	9.311e-04	3.500e-03	0.266	0.790
X30d	5.279e-04	4.867e-03	0.108	0.914
X45d	-2.804e-03	2.983e-03	-0.940	0.347
X90d	5.591e-03	5.494e-03	1.018	0.309
X200d	-2.616e-03	5.139e-03	-0.509	0.611
Mkt..Cap	4.964e-05	1.062e-03	0.047	0.963
X24h.Vol	9.059e-06	1.117e-03	0.008	0.994
Circ..Supply	-1.920e-03	1.103e-03	-1.741	0.082
Total.Supply	9.642e-12	3.189e-13	30.230	<2e-16 ***
Team	1.977e-02	4.204e-02	0.470	0.638
Product	-6.773e-03	1.328e-02	-0.510	0.610
Coin	-8.403e-03	1.685e-02	-0.499	0.618
Social	3.590e-03	2.370e-02	0.151	0.880
Communication	-5.365e-03	8.557e-03	-0.627	0.531
Business	6.493e-03	1.123e-02	0.578	0.563
Avg..volume	2.020e-11	1.844e-09	0.011	0.991
Age..mo.	1.896e-02	1.633e-02	1.161	0.246

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.24 on 962 degrees of freedom  
(40 observations deleted due to missingness)

Multiple R-squared: 0.4896, Adjusted R-squared: 0.4785

F-statistic: 43.95 on 21 and 962 DF, p-value: < 2.2e-16

1. 582.295129943916 2. 2.15332619393653 3. -173.530306079791

In [23]: vif(mod)

Price 1.02423943344359 X1h 1.27695002473311 X24h 2.20857264825821 X7d 2.20262213596122  
X14d 51.0399053125082 X30d 57.8341308908584 X45d 16.3502969281802 X90d 6.77458705855804  
X200d 3.77206853766809 Mkt..Cap 1.02694209743189 X24h.Vol 1.02876014673704 Circ..Supply  
1.07660848685341 Total.Supply 1.01072167422902 Team 6.49637787342385 Product  
1.9710382752697 Coin 1.90590139973695 Social 3.14333047583439 Communication  
1.53408207175482 Business 1.36185221776999 Avg..volume 1.02165670821974 Age..mo.  
1.3028447664378

This is not good, there is some multicollinearity. Let's remove the guilty features.

```
In [26]: mod <- lm(label_growth_rate_Price ~ Price + X1h + X24h + X7d + X30d + X90d + X200d +
  Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team + Product +
  Coin + Social + Communication + Business + Avg..volume + Age..mo.,
  data = price_train_data)
```

summary(mod)

```

pred_lm = predict(mod, newdata = price_test_data)

all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
            price_train_data$label_growth_rate_Price)

```

Call:

```

lm(formula = label_growth_rate_Price ~ Price + X1h + X24h + X7d +
    X30d + X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply +
    Total.Supply + Team + Product + Coin + Social + Communication +
    Business + Avg..volume + Age..mo., data = price_train_data)

```

Residuals:

	Min	1Q	Median	3Q	Max
	-95.222	-0.426	-0.009	0.496	157.685

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-1.506e+00	1.452e+00	-1.037	0.2999
Price	-1.311e-05	9.783e-04	-0.013	0.9893
X1h	2.299e-02	2.857e-02	0.805	0.4212
X24h	-9.303e-03	9.536e-03	-0.976	0.3295
X7d	-3.687e-03	7.124e-03	-0.517	0.6049
X30d	-8.050e-05	9.549e-04	-0.084	0.9328
X90d	1.554e-03	3.466e-03	0.448	0.6541
X200d	-4.926e-04	4.608e-03	-0.107	0.9149
Mkt..Cap	5.801e-05	1.061e-03	0.055	0.9564
X24h.Vol	7.150e-05	1.114e-03	0.064	0.9489
Circ..Supply	-1.902e-03	1.102e-03	-1.726	0.0846 .
Total.Supply	9.631e-12	3.186e-13	30.233	<2e-16 ***
Team	1.702e-02	4.190e-02	0.406	0.6847
Product	-6.461e-03	1.327e-02	-0.487	0.6263
Coin	-7.157e-03	1.676e-02	-0.427	0.6695
Social	5.232e-03	2.359e-02	0.222	0.8245
Communication	-5.266e-03	8.549e-03	-0.616	0.5381
Business	6.734e-03	1.122e-02	0.600	0.5484
Avg..volume	1.029e-11	1.843e-09	0.006	0.9955
Age..mo.	2.052e-02	1.613e-02	1.272	0.2037

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.236 on 964 degrees of freedom

(40 observations deleted due to missingness)

Multiple R-squared: 0.4891, Adjusted R-squared: 0.4791

F-statistic: 48.58 on 19 and 964 DF, p-value: < 2.2e-16

1. 564.062942832866 2. 1.90025594030086 3. -168.065604361777

```
In [27]: vif(mod)
```

```
Price 1.02383458947674 X1h 1.26378621144891 X24h 2.10915904620629 X7d 2.1991343191775
X30d 2.22872902583279 X90d 2.69968306031107 X200d 3.03647959935625 Mkt..Cap
1.02665757634139 X24h.Vol 1.02508903973853 Circ..Supply 1.07610953486792 Total.Supply
1.00943241102529 Team 6.46108652227192 Product 1.96792266463445 Coin 1.88882024934975
Social 3.11853701685625 Communication 1.53321255914753 Business 1.36088711654353
Avg..volume 1.02158918501001 Age..mo. 1.27210148432654
```

This still does not look good. Let's narrow down the features.

```
In [20]: mod <- lm(label_growth_rate_Price ~ Circ..Supply + Total.Supply + 0,
data = price_train_data)
```

```
summary(mod)
```

```
pred_lm = predict(mod, newdata = price_test_data)
```

```
all_metrics(price_test_data$label_growth_rate_Price, pred_lm,
price_train_data$label_growth_rate_Price)
```

Call:

```
lm(formula = label_growth_rate_Price ~ Circ..Supply + Total.Supply +
0, data = price_train_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-96.415	-0.701	-0.407	0.164	158.564

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
Circ..Supply	-2.534e-03	8.533e-04	-2.97	0.00305 **
Total.Supply	9.579e-12	3.097e-13	30.93	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7.074 on 1017 degrees of freedom

(5 observations deleted due to missingness)

Multiple R-squared: 0.485, Adjusted R-squared: 0.484

F-statistic: 478.9 on 2 and 1017 DF, p-value: < 2.2e-16

1. 547.59691212069 2. 1.94101815834009 3. -166.769645220578

We can compare it to the baseline (predicting negative all the time). We assume that we invest the same amount of money in each coin, ignoring fees, and compute the result growth rate of our portfolio following our regression strategy and the baseline.

```
In [29]: filter_vec = !is.na(pred_lm)
          sum((pred_lm*price_test_data$label_growth_rate_Price)[filter_vec])
          /sum(pred_lm[filter_vec])

          # Baseline
          sum(price_test_data$label_growth_rate_Price[filter_vec])
          /length(price_test_data$label_growth_rate_Price[filter_vec])

0.142130885876864
-0.342125271527515
We beat it!
```

### 7.3.3 Portfolio determination in a log normal stock model

#### a) Regression

With our new log-normal assumption we can try to predict the difference of the logarithms of the price between March and December.

```
In [30]: df$price_log_diff = log(df$label_Price) - log(df$Price) # follow  $\mu T + \sigma B_T$ 

split = sample.split(df$price_log_diff, SplitRatio = 0.7)
price_train_data <- filter(df, split== TRUE)
price_test_data <- filter(df, split== FALSE)

price_train_data = price_train_data %>% filter(as.integer(label_disappeared)==1)
# we keep label_disappeared = FALSE
price_test_data = price_test_data %>% filter(as.integer(label_disappeared)==1)

mod_log <- lm(price_log_diff ~ Price + X1h + X24h + X7d + X14d + X30d + X45d + X90d +
              X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team +
              Product + Coin + Social + Communication + Business + Avg..volume + Age..mo.,
              data = price_train_data)

summary(mod_log)

pred_lm = predict(mod_log, newdata = price_test_data)

all_metrics(price_test_data$price_log_diff, pred_lm, price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ Price + X1h + X24h + X7d + X14d +
    X30d + X45d + X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply +
    Total.Supply + Team + Product + Coin + Social + Communication +
    Business + Avg..volume + Age..mo., data = price_train_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-6.2257 -0.4446 0.0628 0.5709 3.8110

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-9.089e-01	1.915e-01	-4.747	2.37e-06 ***
Price	2.427e-05	3.484e-05	0.697	0.4861
X1h	-3.607e-03	5.843e-03	-0.617	0.5372
X24h	-3.514e-03	1.624e-03	-2.164	0.0307 *
X7d	-1.298e-03	1.028e-03	-1.262	0.2074
X14d	-6.971e-04	3.174e-04	-2.196	0.0283 *
X30d	5.989e-04	4.839e-04	1.238	0.2161
X45d	6.502e-05	2.585e-04	0.252	0.8014
X90d	5.886e-04	5.132e-04	1.147	0.2517
X200d	-9.859e-04	5.229e-04	-1.885	0.0597 .
Mkt..Cap	-2.969e-04	1.388e-04	-2.139	0.0327 *
X24h.Vol	-7.968e-07	1.410e-04	-0.006	0.9955
Circ..Supply	2.499e-04	1.455e-04	1.718	0.0862 .
Total.Supply	1.668e-14	1.874e-14	0.890	0.3738
Team	7.936e-03	5.567e-03	1.426	0.1543
Product	-1.076e-03	1.750e-03	-0.615	0.5385
Coin	-2.972e-03	2.195e-03	-1.354	0.1761
Social	-7.665e-04	3.138e-03	-0.244	0.8071
Communication	3.429e-04	1.144e-03	0.300	0.7645
Business	-3.736e-03	1.468e-03	-2.544	0.0111 *
Avg..volume	1.576e-10	1.418e-10	1.112	0.2665
Age..mo.	3.241e-03	2.212e-03	1.465	0.1433

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9755 on 978 degrees of freedom

(37 observations deleted due to missingness)

Multiple R-squared: 0.05583, Adjusted R-squared: 0.03555

F-statistic: 2.754 on 21 and 978 DF, p-value: 3.95e-05

1. 1.37419565150063 2. 0.786600323768397 3. -0.0510270204478771

In [31]: vif(mod\_log)

Price 1.06407441227154 X1h 1.07115978177007 X24h 1.22544188241185 X7d 1.31330264496053  
X14d 753.264205918813 X30d 1730.32601194076 X45d 491.069203199915 X90d 10.6815410522711  
X200d 3.59844428634785 Mkt..Cap 1.0247473315568 X24h.Vol 1.02517852100103 Circ..Supply  
1.07871229383061 Total.Supply 1.01891666805316 Team 6.3719217947138 Product  
1.92537905891746 Coin 1.86682147289798 Social 3.14543604788404 Communication  
1.52506691080304 Business 1.3559062673666 Avg..volume 1.09292020780236 Age..mo.  
1.31230648270897

In [34]: mod\_log <- lm(price\_log\_diff ~ Price + X1h + X24h + X7d + X14d + X90d + X200d +  
Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply + Team + Product +

```

      Coin + Social + Communication + Business + Avg..volume + Age..mo.,
      data = price_train_data)

summary(mod_log)

pred_lm = predict(mod_log, newdata = price_test_data)

all_metrics(price_test_data$price_log_diff, pred_lm, price_train_data$price_log_diff)

Call:
lm(formula = price_log_diff ~ Price + X1h + X24h + X7d + X14d +
    X90d + X200d + Mkt..Cap + X24h.Vol + Circ..Supply + Total.Supply +
    Team + Product + Coin + Social + Communication + Business +
    Avg..volume + Age..mo., data = price_train_data)

Residuals:
    Min       1Q   Median       3Q      Max
-6.1966 -0.4454  0.0553  0.5753  3.7657

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  -9.555e-01  1.909e-01  -5.005 6.62e-07 ***
Price         2.525e-05  3.491e-05   0.723 0.469718
X1h          -3.627e-03  5.851e-03  -0.620 0.535459
X24h          -3.597e-03  1.626e-03  -2.212 0.027225 *
X7d           -1.328e-03  1.029e-03  -1.291 0.196887
X14d          -4.881e-05  2.479e-05  -1.969 0.049266 *
X90d           8.115e-04  4.038e-04   2.010 0.044752 *
X200d         -1.469e-03  4.261e-04  -3.448 0.000589 ***
Mkt..Cap      -2.973e-04  1.391e-04  -2.138 0.032777 *
X24h.Vol      -6.626e-06  1.412e-04  -0.047 0.962579
Circ..Supply   2.468e-04  1.458e-04   1.693 0.090731 .
Total.Supply   1.546e-14  1.874e-14   0.825 0.409431
Team           8.054e-03  5.568e-03   1.446 0.148377
Product       -1.288e-03  1.751e-03  -0.735 0.462238
Coin          -2.875e-03  2.198e-03  -1.308 0.191143
Social        -5.695e-04  3.138e-03  -0.181 0.856039
Communication  4.538e-04  1.145e-03   0.396 0.692061
Business      -3.784e-03  1.471e-03  -2.573 0.010242 *
Avg..volume    1.605e-10  1.420e-10   1.130 0.258656
Age..mo.       3.300e-03  2.204e-03   1.497 0.134725
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9776 on 980 degrees of freedom
(37 observations deleted due to missingness)
Multiple R-squared:  0.04984, Adjusted R-squared:  0.03142

```



F-statistic: 2.706 on 19 and 980 DF, p-value: 0.0001092

1. 1.37718128083273 2. 0.789451576546088 3. -0.0533105214163541

```
In [35]: vif(mod_log)
```

```
Price 1.06389350752336 X1h 1.06954219074384 X24h 1.22325176284725 X7d 1.30815268834064
X14d 4.57576438585392 X90d 6.58402089382897 X200d 2.37953671288416 Mkt..Cap
1.02437219200182 X24h.Vol 1.02414192975333 Circ..Supply 1.07809508758456 Total.Supply
1.01415980805924 Team 6.34760662656953 Product 1.92036647546332 Coin 1.86268778843273
Social 3.13312064680212 Communication 1.52216803443362 Business 1.3550512673174
Avg..volume 1.09262600492207 Age..mo. 1.29729721274557
```

The VIF is okay now. Let's remove some features.

```
In [38]: mod_log <- lm(price_log_diff ~ X24h + X200d + Mkt..Cap + Business,
                        data = price_train_data)
```

```
summary(mod_log)
```

```
pred_lm_log = predict(mod_log, newdata = price_test_data)
```

```
all_metrics(price_test_data$price_log_diff, pred_lm_log,
            price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ X24h + X200d + Mkt..Cap + Business,
    data = price_train_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.2460	-0.4330	0.0580	0.5923	3.9618

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.7949174	0.1167271	-6.810	1.68e-11 ***
X24h	-0.0047998	0.0014715	-3.262	0.00114 **
X200d	-0.0008992	0.0002755	-3.265	0.00113 **
Mkt..Cap	-0.0003394	0.0001378	-2.464	0.01392 *
Business	-0.0027869	0.0012598	-2.212	0.02717 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9794 on 1000 degrees of freedom  
(32 observations deleted due to missingness)

Multiple R-squared: 0.03057, Adjusted R-squared: 0.02669

F-statistic: 7.883 on 4 and 1000 DF, p-value: 2.956e-06

1. 1.35643667936919 2. 0.784982884758303 3. -0.0401929076627052

```
In [39]: filter_vec = !is.na(pred_lm_log)
sum((exp(pred_lm_log)*price_test_data$label_growth_rate_Price)[filter_vec])
      /sum(exp(pred_lm_log)[filter_vec])

# Baseline
sum(price_test_data$label_growth_rate_Price[filter_vec])
      /length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.461436810190696

0.439390893108447

That's fine but not amazing, let's think a bit and try something else.

```
In [40]: mod_log <- lm(price_log_diff ~ X7d + Mkt..Cap + Age..mo., data = price_train_data)

summary(mod_log)

pred_lm_log = predict(mod_log, newdata = price_test_data)

all_metrics(price_test_data$price_log_diff, pred_lm_log,
            price_train_data$price_log_diff)
```

Call:

```
lm(formula = price_log_diff ~ X7d + Mkt..Cap + Age..mo., data = price_train_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.3005	-0.4493	0.0632	0.5849	3.9844

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-1.0571347	0.0505030	-20.932	<2e-16 ***
X7d	-0.0016871	0.0009027	-1.869	0.0619 .
Mkt..Cap	-0.0003295	0.0001397	-2.359	0.0185 *
Age..mo.	0.0041164	0.0019420	2.120	0.0343 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9966 on 1033 degrees of freedom

Multiple R-squared: 0.01299, Adjusted R-squared: 0.01012

F-statistic: 4.531 on 3 and 1033 DF, p-value: 0.003655

1. 1.3321906348193 2. 0.766205419022985 3. 0.0152780706337514

```
In [41]: filter_vec = !is.na(pred_lm_log)
          sum((exp(pred_lm_log)*price_test_data$label_growth_rate_Price)[filter_vec])
          /sum(exp(pred_lm_log)[filter_vec])

# Baseline
          sum(price_test_data$label_growth_rate_Price[filter_vec])
          /length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.500360688948723

0.412744498999508

Much better already! We see that filtering out coins where informations are missing, our portfolio behaves already better! That's because trustworthy coins get all their information filled.

#### b) Random forest

We take the features from the 2 previous models and we try it out on a random forest.

```
In [50]: price_test_data_filled_with_0 <- price_test_data
          price_test_data_filled_with_0[is.na(price_test_data_filled_with_0)] <- 0
          price_test_data_log.mm = as.data.frame(model.matrix(price_log_diff ~ X24h + X7d +
            X200d + Mkt..Cap + Age..mo. + Business, data = price_test_data_filled_with_0))

          train.rf <- train(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. +
            Business,
            data = price_train_data,
            method = "rf",
            na.action = na.omit,
            tuneGrid = data.frame(mtry=1:6),
            trControl = trainControl(method="cv", number=5),
            distribution="gaussian",
            metric = "RMSE")

          train.rf$results
          train.rf
          best.rf <- train.rf$finalModel
          pred.rf <- predict(best.rf, newdata = price_test_data_log.mm)
          all_metrics(price_test_data$price_log_diff, pred_lm_log,
            price_train_data$price_log_diff)
```

mtry <int>	RMSE <dbl>	Rsquared <dbl>	MAE <dbl>	RMSESD <dbl>	RsquaredSD <dbl>	MAESD <dbl>
1	0.9868626	0.02396270	0.7021205	0.09754761	0.03461219	0.04148504
2	1.0030523	0.02013034	0.7180743	0.09988699	0.03206543	0.03882143
3	1.0138758	0.01549198	0.7273112	0.10018572	0.02833747	0.03858953
4	1.0172491	0.01483976	0.7298685	0.09965039	0.02697329	0.04055537
5	1.0190250	0.01714539	0.7307562	0.10116483	0.03156220	0.03939837
6	1.0224764	0.01754349	0.7330653	0.10168716	0.03172612	0.04190727

Random Forest

1037 samples  
6 predictor

No pre-processing

Resampling: Cross-Validated (5 fold)

Summary of sample sizes: 804, 805, 805, 804, 802

Resampling results across tuning parameters:

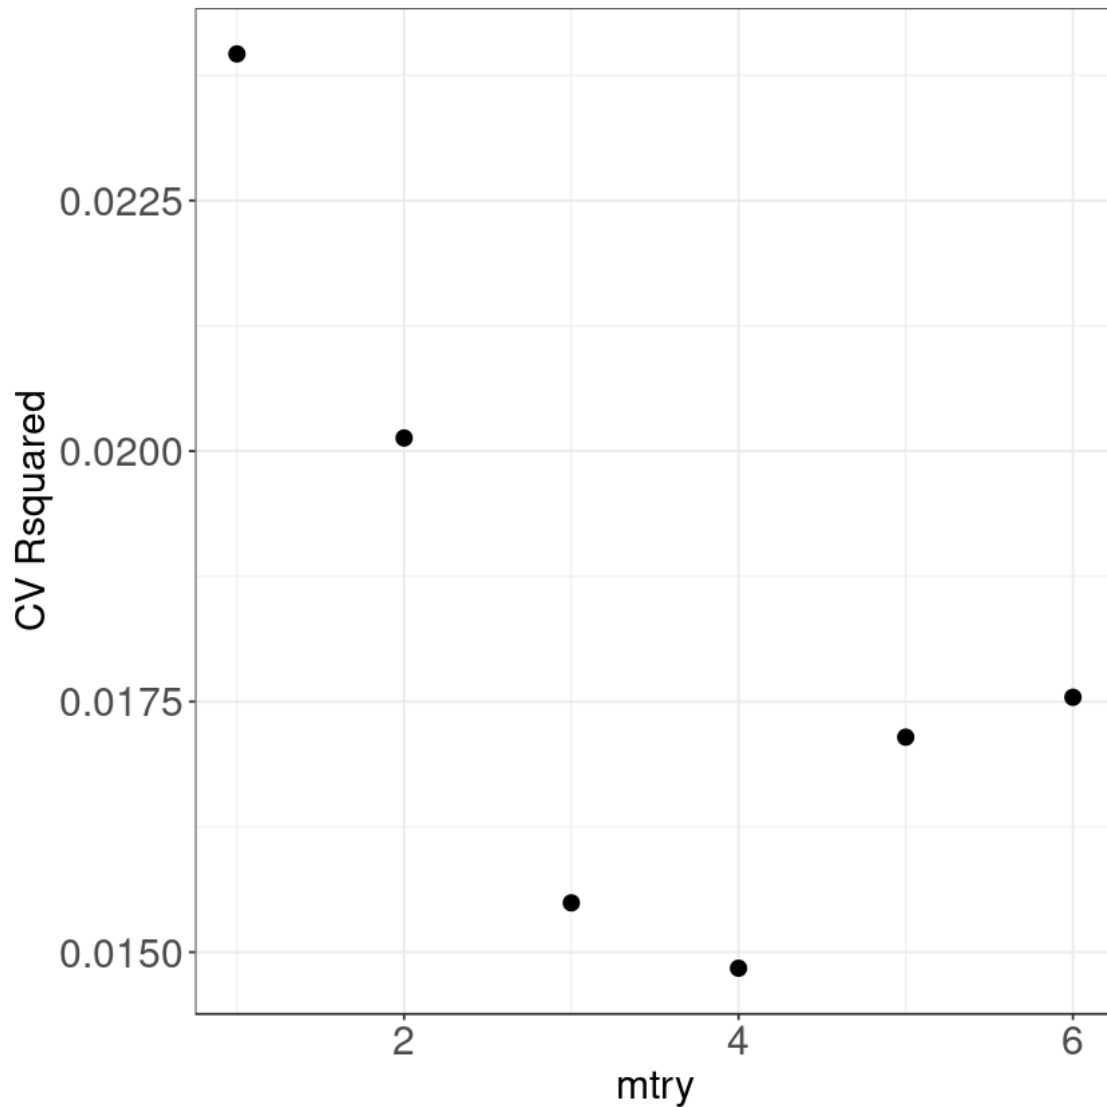
mtry	RMSE	Rsquared	MAE
1	0.9868626	0.02396270	0.7021205
2	1.0030523	0.02013034	0.7180743
3	1.0138758	0.01549198	0.7273112
4	1.0172491	0.01483976	0.7298685
5	1.0190250	0.01714539	0.7307562
6	1.0224764	0.01754349	0.7330653

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was mtry = 1.

1. 1.3321906348193 2. 0.766205419022985 3. 0.0152780706337514

```
In [51]: ggplot(train.rf$results, aes(x = mtry, y = Rsquared)) + geom_point(size = 3) +  
         ylab("CV Rsquared") + theme_bw() + theme(axis.title=element_text(size=18),  
         axis.text=element_text(size=18))
```



The best value is 1 for mtry.

```
In [52]: filter_vec = !is.na(pred.rf)
          sum((exp(pred.rf)*price_test_data$label_growth_rate_Price)[filter_vec])
          /sum(exp(pred.rf)[filter_vec])

          # Baseline
          sum(price_test_data$label_growth_rate_Price[filter_vec])
          /length(price_test_data$label_growth_rate_Price[filter_vec])
```

0.318126401502254

0.410949916909122

Unfortunately this model does not beat the baseline.

c) Boosting

We use the same features.

```
In [54]: tGrid = expand.grid(n.trees = (1:10)*1000, interaction.depth = c(1,2,4,6,8,10),
                             shrinkage = 0.001, n.minobsinnode = 10)

set.seed(42)
train.boost <- train(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. +
                     Business,
                     data = price_train_data,
                     method = "gbm",    ## gradient boosting machine
                     tuneGrid = tGrid,
                     trControl = trainControl(method="cv", number=5),
                     distribution = "gaussian",
                     metric = "RMSE",
                     na.action = na.omit)

train.boost
best.boost <- train.boost$finalModel

ggplot(train.boost$results, aes(x = n.trees, y = Rsquared, colour =
  as.factor(interaction.depth))) + geom_line() +
  ylab("CV Rsquared") + theme_bw() + theme(axis.title=element_text(size=18),
  axis.text=element_text(size=18)) +
  scale_color_discrete(name = "interaction.depth")
```

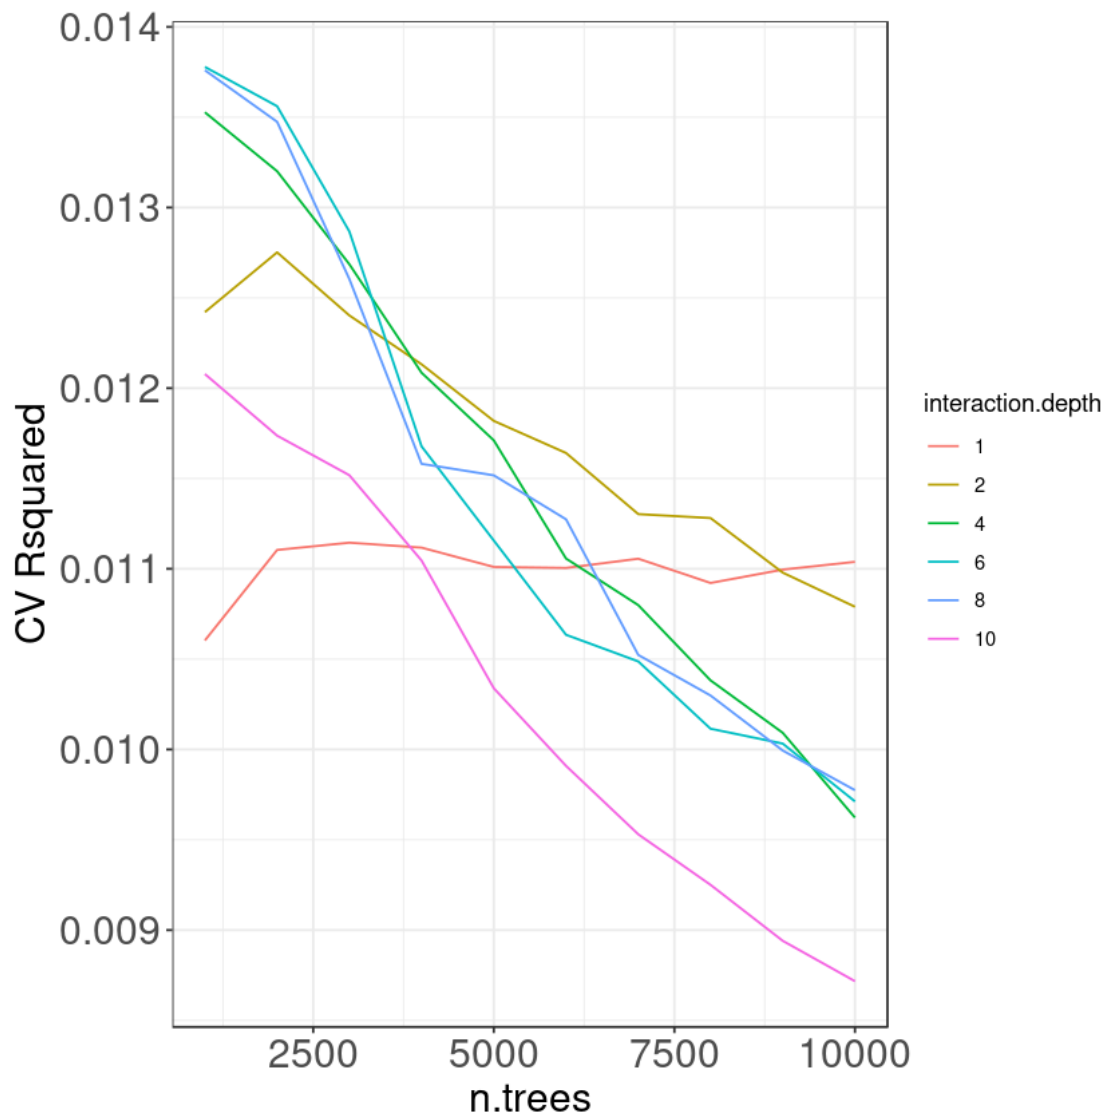
Stochastic Gradient Boosting

1037 samples  
6 predictor

No pre-processing  
Resampling: Cross-Validated (5 fold)  
Summary of sample sizes: 803, 804, 805, 804, 804

Tuning parameter 'shrinkage' was held constant at a value of 0.001

Tuning parameter 'n.minobsinnode' was held constant at a value of 10  
RMSE was used to select the optimal model using the smallest value.  
The final values used for the model were n.trees = 1000, interaction.depth =  
6, shrinkage = 0.001 and n.minobsinnode = 10.



The best parameters are the biggest ones.

```
In [56]: pred.best.boost <- predict(best.boost, newdata = price_test_data_log.mm,
  n.trees = 10000, interaction.depth=10) # from CV plot
```

Warning message in predict.gbm(best.boost, newdata = price\_test\_data\_log.mm, n.trees = 10000, :  
 “Number of trees not specified or exceeded number fit so far. Using 1000.”

It does not want to do it so we need to do it ourselves.

```
In [63]: library(gbm)
```

```
mod.boost <- gbm(price_log_diff ~ X24h + X7d + X200d + Mkt..Cap + Age..mo. + Business,
  data = price_train_data,
```

```

distribution = "gaussian",
n.trees = 10000,
shrinkage = 0.001,
interaction.depth = 10)

# NOTE: we need to specify number of trees to get a prediction for boosting
pred.boost <- predict(mod.boost, newdata = price_test_data_log.mm, n.trees=10000)

In [64]: filter_vec = !is.na(pred.boost)
sum((exp(pred.boost)*price_test_data$label_growth_rate_Price)[filter_vec])
      /sum(exp(pred.boost)[filter_vec])

# Baseline
sum(price_test_data$label_growth_rate_Price[filter_vec])
      /length(price_test_data$label_growth_rate_Price[filter_vec])

0.290123603135029
0.410949916909122
This is also very bad compared to the baseline. Too much overfitting!

```